

Data Intensive Systems (DIS)

KBH-SW7 E25


4. MapReduce

Agenda

- ⌚ **Background**
 - ⌚ Scalability in computation
 - ⌚ Big ideas behind MapReduce
- ⌚ Programming model
- ⌚ Execution framework
- ⌚ Software implementation
- ⌚ MapReduce and databases

Scale-up vs. Scale-out

- ⦿ Demand for efficiency is the drive



Scale-up vs. Scale-out

- Scalability is a system's ability to swiftly enlarge or reduce the power or size of computing, storage, or networking infrastructure.
- In a data system, scalability means its ability to process larger amounts of data and maybe also at a faster pace.
- To achieve scalability, we need to add capacity to the infrastructure.
- **Scale-up:** adding more (virtual) resources, e.g., equipping a computer with more/better CPUs, more RAMs, and larger hard disks.
- **Scale-out:** adding more computers to spread the workload across more machines. This form *clusters* of machine.

Scale-up (Vertical Scaling)

- When to scale up
 - When there's a performance impact, e.g., limited I/O and CPU capacity increase latency and cause performance bottleneck.
 - When storage optimization does not work, i.e., no room for better on the current resources.
- Pros
 - Relative speed up. E.g., a dual processor, a faster DRAM
 - Simplicity. System architecture remains the same.
 - Cost-effectiveness: cheaper than scale-out with many machines
 - Limited energy consumption
- Cons
 - Latency
 - Labor and risks
 - Aging hardware

Scale-out (Horizontal Scaling)

➤ When to scale out

- When a long-term scaling strategy is needed.
- When upgrades need to be flexible.
- When storage workloads need to be distributed.

➤ Pros

- To use newer server technologies
- Adaptability to demand changes
- Cost management (incremental, predictable):
 - We can use commodity low-end servers. This is significantly cheaper than buying high-end servers.

➤ Cons

- Limited rack space
- Increased operational costs
- Higher upfront costs



A cluster of
low-end machines

Programming

- Programming for parallel & distributed systems is very difficult
 - Concurrency control: Race conditions, deadlocks, etc.
 - A programmer needs to spend considerable efforts on ‘infrastructural details’ than on the real problem.
- Ideally
 - Programmers should be liberated from such (system-level) infrastructural details shared by a large class of, if not all, problems.
 - A solution should scale with input data.
 - It should take twice long, if the data size doubles.
 - It should take the same time on a twice big cluster, if the data size doubles.

Big Ideas behind MapReduce

- Scale out, not up
 - Make use of large number of commodity, low-end machines
- Fault-tolerance
 - Assume failures are common
- Move processing to the data
 - Not the other way around, which is time-consuming
- Process data *sequentially* and avoid random access
 - Suitable for some particular types of applications, not all
- Hide system-level details from the application developers
- Seamless scalability

3 Concepts of MapReduce

- The term of MapReduce can refer to three distinct but related concepts.
- Programming model
 - How are developers supposed to program?
- Execution framework
 - Runtime that coordinates the execution of programs written in the MapReduce paradigm
- Software implementation
 - The implementation of the previous two: Google, Hadoop and others.

Agenda

- ⌚ Background
- ⌚ Programming model
- ⌚ Execution framework
- ⌚ Software implementation
- ⌚ MapReduce and databases

Functional Programming Roots


- ➊ Higher-order functions, functions that accept other functions as arguments
 - ➋ Two common built-in functions *map* and *fold*
- ➋ Consider a 4d vector $\mathbf{v}(v_1, v_2, v_3, v_4)$. To get $|\mathbf{v}|$, we need to compute $\sum v_i^2$.
 - ➌ **Transformation:** *map* applies $f(x)=x^2$ to each v_i .
 - ➍ **Aggregation:** *fold* applies $g(x_1, x_2)=x_1+x_2$ to v_i^2 s.

map phase

Transformation can always be parallelized.

reduce phase

Aggregation may enjoy group-based parallelism.



MapReduce Programming Model

- **Input:** A set of key/value pairs
- **Output:** Another set of key/value pairs
- Keys and values can be *primitives* or *complex types*
- A programmer implements two functions: **map** and **reduce**
 - **map:** $(k_1, v_1) \rightarrow \text{list}(k_2, v_2)$
 - Takes an input pair and produces a set of intermediate key/value pairs.
 - MapReduce groups all intermediate pairs with the same key and gives them to *reduce*.
 - **reduce:** $(k_2, \text{list}(v_2)) \rightarrow \text{list}(k_3, v_3)$ (**Hadoop**)
 $\text{list}(v_2)$ (**Google**)
 - Takes an intermediate key and the set of all values for that key.
 - Merges the values to form a smaller set (typically empty or with a single value)


Example: WordCount

- Data: A large collection of documents
- Expected output: #occurrence of each word across *all* documents.
- Parallelism could be easily achieved.
 - A number of mappers can work on different documents in parallel.
 - For each word they see, they output an intermediate count (of 1) for the word.
 - When the mappers all finish, a number of reducers can work in parallel.
 - A reducer gets all counts for a certain word and ‘reduces’ the list of counts (i.e., 1’s) to a single value.

Map and Reduce for WordCount

```
map(String key, String value) :  
    // key: document id; value: document contents  
    for each word w in value:  
        emit(w, count)  
  
reduce(String key, Iterator values) :  
    // key: a word; values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += v;  
    emit(result);
```

A Simplified Conceptual View



Other Examples

- Counting URL accesses
 - Very similar to the WordCount example
 - The map function processes log files and outputs (URL, 1) for each access to URL
 - The reduce function adds all intermediate 1's for each URL
- Reverse web-link graph
 - The map function is given a document's URL as key, and the content as value. For each referenced target, it outputs (target, URL)
 - The reduce functions just outputs the list of URLs that reference a given target

Demo in Python

- Problem
 - Given a large list of strings, return the longest string
- StringLengthMR.py
 - This file implements four versions of MapReduce based functions
- DIS-E23-Lecture4_MapReduce_strings.ipynb
 - This is the ‘main’ file that calls these functions
- This format of modularity is needed to use multiprocessing in Jupyter Notebook
 - To ‘simulate’ or utilize multiple processors for MapReduce



Agenda

- ⌚ Background
- ⌚ Programming model
- ⌚ **Execution framework**
- ⌚ Software implementation
- ⌚ MapReduce and databases

Execution Framework

- MapReduce separates the *what* of distributed processing from the *how*.
- A MapReduce program (a job) consists of
 - Code for **mappers** and **reducers**
 - And **combiners** and **partitioners** to be covered shortly
 - Configuration parameters
 - E.g., where the input is and where the output should go
- The developer submits the job to the *submission node* of a cluster
 - It's called **jobtracker** in Hadoop
- **Execution framework** (a.k.a. runtime) takes care of everything else *transparently*:
 - All other aspects of distributed code execution
 - On clusters ranging from a single node to thousands of nodes

MapReduce Runtime Responsibilities

⦿ Scheduling

- ⦿ Each job is divided into smaller units called tasks.
 - › For map: splits of input key-value pairs
 - › For reducer: division of the intermediate key space.
- ⦿ Tasks are assigned to nodes in the cluster.
- ⦿ Coordination may be needed among tasks and/or nodes.

⦿ Data/code co-location

- ⦿ For data locality, the scheduler starts tasks on the node having the needed data in its local storage.
- ⦿ If it's not possible, new tasks will be started elsewhere with the data streamed into via the network.
Intra-rack transmission is preferred over inter-rack transmission.

MapReduce Runtime Responsibilities, cont.

➤ Synchronization

- ‘Shuffle and sort’ is a barrier between the map and reduce phases.
- It involves copying immediate data over the network.
- A job with m mappers and r reducers involves up to $m \times r$ distinct copy operations.
- Unlike the *fold* operation, the reduce computation cannot start until
 1. all the mappers have finished emitting key-value pairs *and*
 2. all intermediate key-value pairs have been shuffled and sorted.


➤ Error and fault handling

- All responsibilities above are in an error- and fault-prone environment.
 - Think of hardware failures of low-end commodity servers.
 - Software bugs and data errors

Combiners and Partitioners

- Combiners are an optimization that allows for *local aggregation* before Shuffle and Sort.
 - This helps reduce the number of intermediate key-value pairs.
 - Can be regarded as ‘mini-reducers’.
 - Recall the WordCount example
- Partitioners divide up the intermediate key space and assign intermediate key-value pairs to reducers.
 - A partitioner specifies the task (and thus the node to be determined by the runtime) to which an intermediate key-value pair must be copied.
 - The information is for the Shuffle and Sort phase to copy the data to the right place.
 - Hashing is the default partitioning function on the intermediate keys.

The Conceptual View Again



Agenda

- ⌚ Background
- ⌚ Programming model
- ⌚ Execution framework
- ⌚ **Software implementation**
- ⌚ MapReduce and databases

The Distributed File System (DFS)

- In traditional clusters, computation and storage are separated as distinct components.
- DFS abandons the separation and makes data considerably closer to computation, which helps reduce overall processing time.
 - DFS divides user data into blocks and replicates them across the local disks of nodes in the cluster.
 - DFS uses significantly large block sizes.
 - DFS adopts a **master-slave** architecture
 - The master maintains the file namespace
 - The slaves manage the actual data blocks.
- MapReduce can work without DFS, but many advantages will disappear without an underlying DFS or the like.
 - Google File System (GFS)
 - Hadoop Distributed File System (HDFS)

HDFS Architecture

- Master
 - Namenode
- Slaves
 - Datanodes




Figure from S. Sudarshan, IIT Bombay

More Details

- By default, HDFS stores *three* separate copies of each data block to ensure reliability, availability and performance.
 - In large clusters, the three replicas are spread across different physical racks.
- HDFS files are immutable---they only accept appends.
 - WORM (Write Once Read Many)
- HDFS namenode's responsibility
 - Namespace management
 - Coordination file operations
 - Maintaining overall health of the file system
- The namenode forms a single point of failure. So, a warm standby namenode is often used in case that the primary namenode fails.

Concurrent Clients of HDFS

- Read from a replica
- Write to all replicas





Figure from S. Sudarshan, IIT Bombay

Assumptions behind GFS and HDFS

- ⦿ The file system stores a relatively modest number of large files.
 - ⦿ Large multi-block files are favored than too many small files
- ⦿ Workloads are *batch* oriented, dominated by long streaming reads and large sequential writes.
 - ⦿ Batch operations on large chunks of data without data caching.
- ⦿ Applications are aware of the characteristics of the distributed files system.
 - ⦿ Data management is pushed onto the end application.
- ⦿ The file system is deployed in an environment of cooperative users.
 - ⦿ Security is essentially not a concern in MapReduce.
- ⦿ The system is built from unreliable but inexpensive commodity machines.
 - ⦿ Self-monitoring and self-healing mechanisms are needed.

Hadoop Cluster Architecture

- › 3 types of nodes
- › 2 types of trackers
- › 2 types of daemons



The two nodes can be co-located if the cluster is small.

MapReduce on Hadoop Cluster

- A simplified view




Figure from Henrik Bulskov

Execution Overview

- This is about Google MapReduce
- But the overall procedure is similar in Hadoop

The files can be local or distributed




Figure from Jeffrey Dean & Sanjay Ghemawat


Agenda

- ⌚ Background
- ⌚ Programming model
- ⌚ Execution framework
- ⌚ Software implementation
- ⌚ MapReduce and databases

Hadoop vs. Databases

- First of all, Hadoop is *not* a (relational) database system.
- Structured vs. Unstructured Data
 - Hadoop can handle both and is a perfect match for unstructured data.
 - Relational databases are designed for structured data only.
- Scalable analytics infrastructure
 - Constant and predictable workloads are good for databases.
 - Increasing data demands can take advantage of Hadoop.
- Cost-effective
 - Low-end commodity machines and open-source software for Hadoop
- Fast data analysis
 - Databases can do better for time-sensitive and general data analysis.
- It's possible/interesting to build hybrid systems with both Hadoop and databases.

Hadoop Ecosystem



MapReduce vs. Parallel Databases

⌚ Schema

- ⌚ The parallel RDBMS forces data into a schema (structured)
 - › Parsing is only needed at load time and the system enforces constraints
 - › Can improve compression and makes it easy to extract a given attribute
 - › Helps when optimizing (the declarative) queries

⌚ MapReduce

- › No schema – very flexible (unstructured)
- › But the programmer must always parse the input (and check that the constraints are not violated). This takes time...

⌚ Indexing

- ⌚ The parallel RDBMS has indexes.
- ⌚ MapReduce does not have built-in indexes.

MapReduce vs. Parallel Databases, cont.

- Programming model
 - SQL: Declarative programming. You specify *what* you want.
 - MapReduce: Overall declarative; but imperative inside the map/reduce functions.
- Getting started is easier in MapReduce. Maintenance is harder.
- Fault tolerance
 - MapReduce handles it if a node fails and only re-computes the lost part.
 - But MapReduce is typically slower and needs more nodes than an RDBMS. The risk of a failure is then bigger.
 - An RDBMS must restart the entire query execution.

Summary

- ⌚ ‘MapReduce is the most successful abstraction over large-scale computational resources we have seen to date.’ – Lin and Dyer
- ⌚ Like all other abstractions in Computer Science, MapReduce is imperfect
 - ⌚ It makes certain large-scale data processing problems easier, but others still (even more) difficult
- ⌚ In the data world, there is no ‘one-size-fits-all’ solution. One should choose between MapReduce and (parallel) databases according to the characteristics of her own data and workloads.

References

⦿ Mandatory reading

- ⦿ Jimmy Lin and Chris Dyer: Data-Intensive Text Processing with MapReduce. Morgan & Claypool Publishers, 2010. Chapter 2 *MapReduce Basics*.

⦿ Further readings

- ⦿ Jeffrey Dean and Sanjay Ghemawat: MapReduce: Simplified Data Processing on Large Clusters. OSDI 2004: 137-150
- ⦿ Andrew Pavlo, Erik Paulson, Alexander Rasin, Daniel J. Abadi, David J. DeWitt, Samuel Madden, Michael Stonebraker: A comparison of approaches to large-scale data analysis. SIGMOD Conference 2009: 165-178

⦿ Acknowledgements

- ⦿ In addition to the authors above, parts of the slides are also adapted from or inspired by:
 - › Christian Thomsen (AAU), S. Sudarshan (IIT Bombay), Henrik Bulskov (RUC)

Exercises

Write your code using MapReduce in Python and Jupyter Notebook to resolve the following two tasks

1. To count *each* hashtag that appears in the ID-Hashtag file (in Moodle).
2. To return the hashtag with the highest appearance frequency.

NB: Think how these two tasks are correlated but different, and how you may reuse your code as much as possible.