



Data Intensive Systems (DIS)

KBH-SW7 E25

9. Association Rules

A Real Application of Association Rules

- Amazon's recommendation
 - 90% buyers who bought A also bought B.
 - Since you've bought A, you may also want B.

Make recommendations based on rules of high support, confidence and lift.

Frequently Bought Together

Price for all three: \$67.41

Add all three to Cart

Add all three to Wish List

Show availability and shipping details

This item: Big Data: A Revolution That Will Transform How We Live, Work, and Think by Viktor Mayer-Schönberger Paperback \$10.61

Data Science for Business: What you need to know about data mining and data-analytic thinking by Foster Provost Paperback \$37.99

Predictive Analytics: The Power to Predict Who Will Click, Buy, Lie, or Die by Eric Siegel Hardcover \$18.81

Have one to sell? Sell on Amazon

Customers Who Bought This Item Also Bought

Item	Author	Price	Rating	Reviews
Data Science for Business: What you need...	Foster Provost	\$37.99	4.5	103
Predictive Analytics: The Power to Predict Who...	Eric Siegel	\$18.81	4.5	235
Big Data: Understanding How Data Powers Big...	Bill Schmarzo	\$28.94	4.5	5
Big Data For Dummies	Judith Hurwitz	\$19.81	4.5	31
Digital Sociology	Deborah Lupton	\$46.95	4.5	Prime
Big Data: Using SMART Big Data, Analytics...	Bernard Marr	\$16.74	4.5	9

Page 2

Agenda

- ⌚ Problem definition
 - ⌚ Support, confidence, lift, and association rule
 - ⌚ Frequent itemsets
 - ⌚ Steps for association rule mining
- ⌚ Apriori principle, Apriori algorithm
- ⌚ Deriving association rules from frequent itemsets
- ⌚ Reflection on Apriori algorithm

Market Basket Data

- Large set of *items*, i.e., things sold in a supermarket
- Large set of *baskets*, each a small subset of items, i.e., things that one customer buys in one **transaction**
- Transaction table T: market-basket data
 - Each record is a transaction, containing a set of items
 - Many-to-many mapping (association) between items and baskets
- What can we do with this type of data?
 - E.g., counting whether the combination {Milk, Bread} is **frequent** or not

TID	Items
1	{Milk, Bread, Beer, Diapers}
2	{Bread, Eggs}
3	{Bread, Diapers}
4	{Milk, Bread, Cola}
5	{Milk, Bread, Diapers}

Transaction table

What Is Association Rule Mining?

- Finding **frequent patterns** and **associations** (rules) among sets of items in a transaction table
- Motivation (market basket analysis):
 - How likely is that the customers buying *milk* are also buying *bread*?
 - Such rules help retailers making decisions
 - Plan the shelf space: placing milk close to bread, more convenient for the customers
 - Offer promotions/discounts for those products together

What Is an Association Rule?

- ⌚ An **association rule** correlates (associates) the presence of one set of items with that of another set of items
- ⌚ Examples
 - ⌚ Rule form: **Body \Rightarrow Head [support, confidence]**
 - ⌚ **milk \Rightarrow bread [5%, 70%]**
 - › 5% of transactions buy both milk and bread
 - › transactions that buy milk have 70% chance of buying also bread
- ⌚ Applications: basket data analysis, catalog design
 - ⌚ * \Rightarrow **chocolate** (How to boost the sales of chocolate)
 - ⌚ **Home Electronics \Rightarrow *** (What other products should the store stock up?)

Rule Components

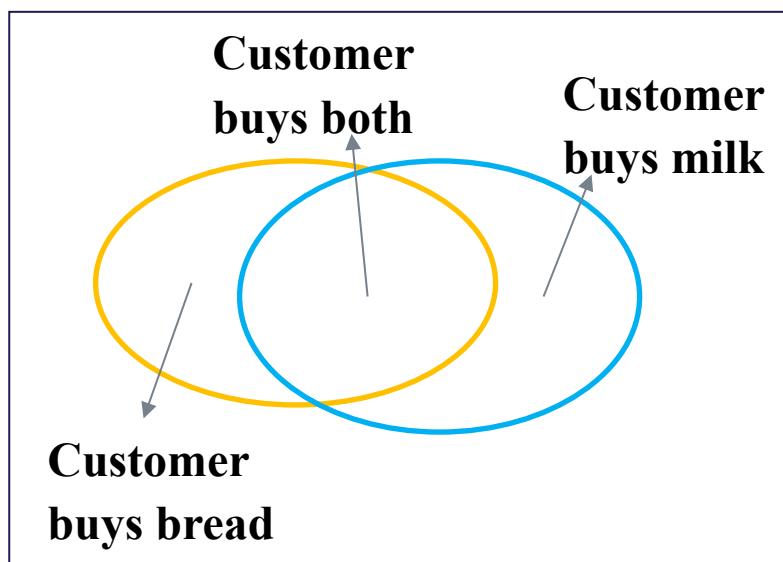
- ⦿ An **itemset** means a set of items
 - ⦿ E.g., {a, b, c}
- ⦿ Let T be a collection of transactions
 - ⦿ E.g., T = {TID 1, TID 2, TID 3, TID 4}
- ⦿ Let I be the set of items that appear in the database,
e.g., I = {a, b, c, g, e, f}
- ⦿ A rule is defined by $X \Rightarrow Y$, where $X \subset I$, $Y \subset I$, and $X \cap Y = \emptyset$
 - ⦿ E.g., {b,c} \Rightarrow {e} is a rule. We may simplify it as bc \Rightarrow e if the context is clear.
 - ⦿ E.g., {b,c} \Rightarrow {c, e} is not a rule

TID	Items
1	a,b,c
2	a,c
3	a,g
4	b,e,f

Interesting Rules

- A rule is said to be **interesting** (or valid) when:
 - Its items appear frequently in the transaction table (**support**)
 - It holds with a high probability (**confidence**)

Example: $milk \Rightarrow bread$



NB: X and Y are itemsets.

Find all the rules $X \Rightarrow Y$ with confidence and support above given thresholds

- **support s**, probability that a transaction contains $X \cup Y$
- **confidence c**, conditional probability that a transaction having X also contains Y

Example (1)

- Find the support and confidence of the rule: $\{B,D\} \Rightarrow \{A\}$

- Support value of $sup(ABD)$:

- percentage of tuples with $\{A,B,D\}$
 $= (3/4)*100\% = 75\%$

- Confidence value of $conf(BD \Rightarrow A)$

$$\frac{\text{number of transactions that contain } \{A, B, D\}}{\text{number of transactions that contain } \{B, D\}} = \frac{3}{3} = 100\%$$

TID	items bought
100	{F,A,D,B}
200	{D,A,C,E,B}
300	{C,A,B,E}
400	{B,A,D}

$$prob(Y | X) = \frac{prob(X \cup Y)}{prob(X)}$$

$$conf(X \Rightarrow Y) = \frac{sup(X \cup Y)}{sup(X)} = \frac{frequency(X \cup Y)}{frequency(X)}$$

Example (1)

- Find interesting rules

Transaction ID	Items Bought
2000	A,B,C
1000	A,C
4000	A,D
5000	B,E,F

Thresholds:

Min. support 50%
Min. confidence 50%

Frequent Itemset	Support
{A}	75%
{B}	50%
{C}	50%
{A,C}	50%

$A \Rightarrow C$ is a valid rule because:

$$\text{support} = \text{support}(\{A \cup C\}) = 2/4 = 50\%$$

$$\text{confidence} = \text{support}(\{A \cup C\})/\text{support}(\{A\}) = 50\%/75\% = 66.6\%$$

Lift of A Rule

- ⦿ $\text{Lift}(X \Rightarrow Y) = \text{confidence}(X \Rightarrow Y)/\text{support}(Y)$
= $\text{support}(X \cup Y) / (\text{support}(X) * \text{support}(Y))$
= $(\text{frequency}(X \cup Y) * |T|) / (\text{frequency}(X) * \text{frequency}(Y))$

- ⦿ Lift($X \Rightarrow Y$) refers to the increase in the ratio of sale of Y when X is sold
 - ⦿ Lift = 1: No association between products X and Y.
 - ⦿ Lift > 1: Products X and Y are more likely to be bought together.
 - ⦿ Lift < 1: The two products are unlikely to be bought together.

Example of Lift

Transaction ID	Items Bought
2000	A,B,C
1000	A,C
4000	A,D
5000	B,E,F

Frequent Itemset	Support
{A}	75%
{B}	50%
{C}	50%
{A,C}	50%

Rule $A \Rightarrow C$:

$$\text{support} = \text{support}(\{A \cup C\}) = 2/4 = 50\%$$

$$\text{confidence} = \text{support}(\{A \cup C\})/\text{support}(\{A\}) = 50\%/75\% = 66.6\%$$

$$\text{lift} = \text{confidence}(A \Rightarrow C)/\text{support}(C) = 66.6\%/50\% = 1.333$$

Lift's meaning: the likelihood of buying a A and C *together* is 1.33 times more than the likelihood of just buying the C.

Recommendation in Amazon

- Two types of recommendation
 - 90% buyers who bought A also bought B.
 - Since you've bought A, you may also want B.

The screenshot shows an Amazon product page for the book "Big Data: A Revolution That Will Transform How We Live, Work, and Think" by Viktor Mayer-Schönberger. The page features two main recommendation sections:

- Frequent itemset:** This section is highlighted with a green box and labeled "Frequent itemset". It shows three books that are frequently bought together: "BIG DATA", "Data Science for Business", and "Predictive Analytics". Buttons for "Add all three to Cart" and "Add all three to Wish List" are present. Below the books, a list of items includes:
 - This item: Big Data: A Revolution That Will Transform How We Live, Work, and Think by Viktor Mayer-Schönberger Paperback \$10.61
 - Data Science for Business: What you need to know about data mining and data-analytic thinking by Foster Provost Paperback \$37.99
 - Predictive Analytics: The Power to Predict Who Will Click, Buy, Lie, or Die by Eric Siegel Hardcover \$18.81
- Association rules:** This section is highlighted with a green box and labeled "Association rules". It shows a list of books that customers who bought the main item also bought:
 - Data Science for Business: What you need to know about data mining and data-analytic thinking by Foster Provost Paperback \$37.99
 - Predictive Analytics: The Power to Predict Who Will Click, Buy, Lie, or Die by Eric Siegel Hardcover \$18.81
 - Big Data: Understanding How Data Powers Big... by Bill Schmarzo Paperback \$28.94 Prime
 - Big Data For Dummies by Judith Hurwitz Paperback \$19.81 Prime
 - Digital Sociology by Deborah Lupton Paperback \$46.95 Prime
 - Big Data: Using SMART Big Data, Analytics... by Bernard Marr Paperback \$16.74 Prime

Other visible elements include "Add to Wish List", "Have one to sell?", "Sell on Amazon", and a "Page" button.

Causality vs. Correlation

- Causality
 - From the very first day, humans are curious about *why*.
 - With big data, it may be very hard to see the exact reasons.
- Correlation
 - Instead, we can find interesting patterns or associations of different things from big data.
 - Probability instead of certainty (not totally random).
 - Association rule mining.
- **NB:** Association rules are empiricism! What they tell may not be the true cause and effect.

Steps of Association Rule Mining

1. Find the *frequent itemsets*

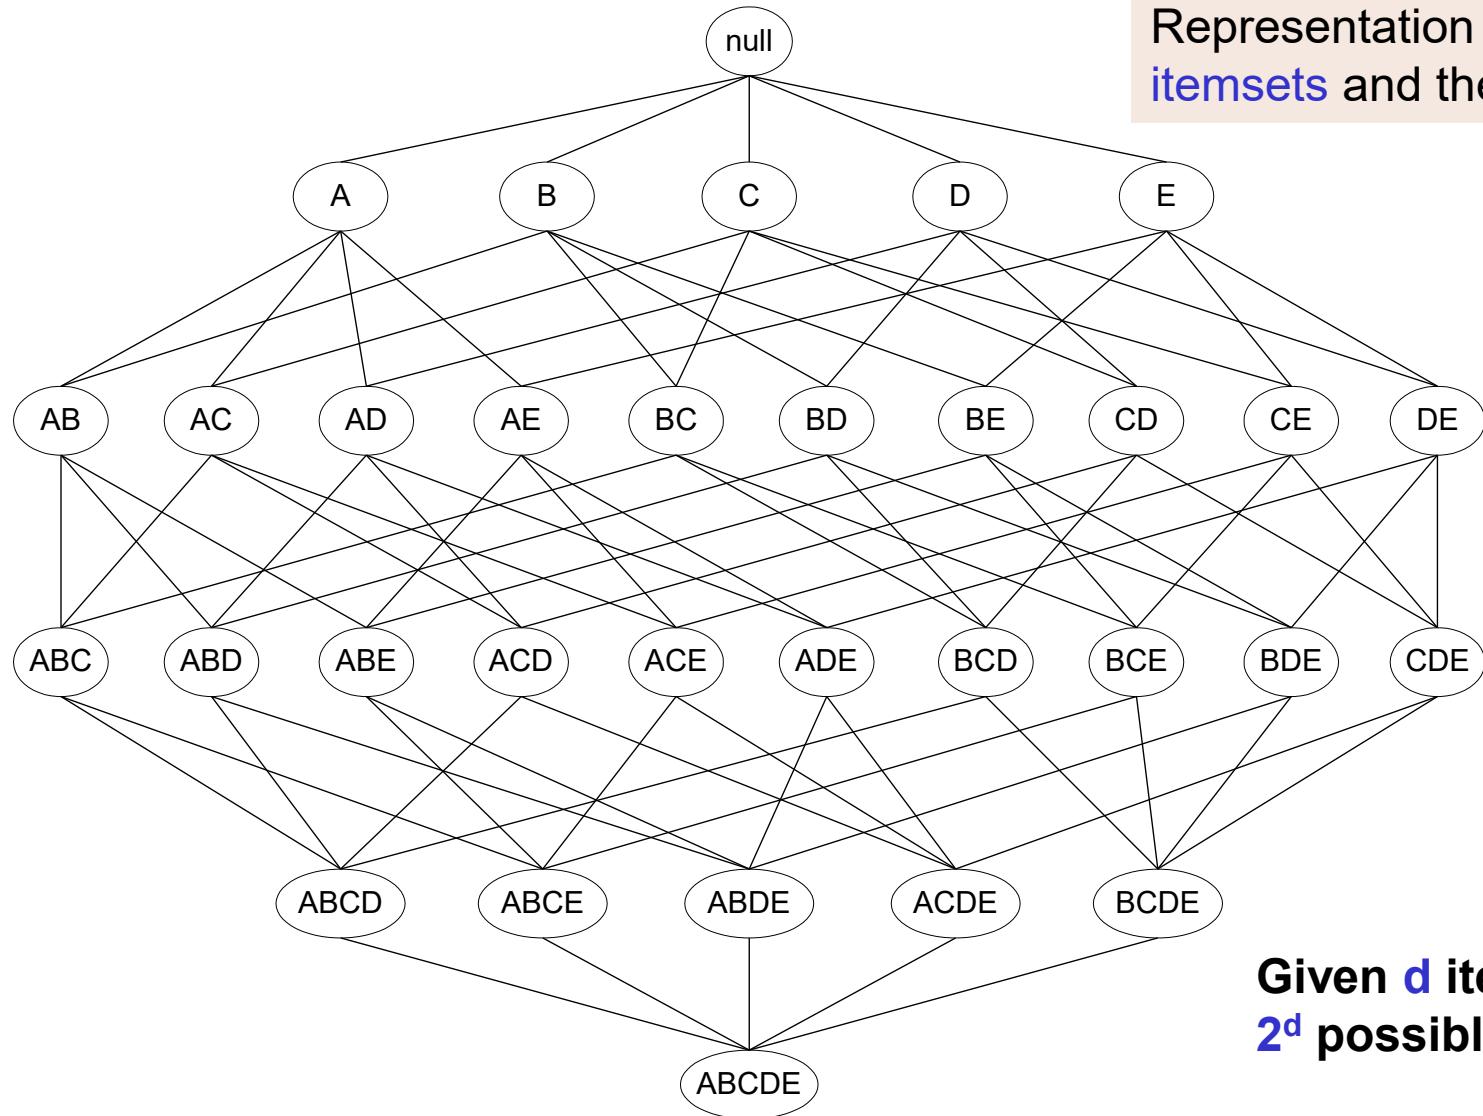
- The sets of items that have minimum support
- How to do this efficiently?

2. Use the frequent itemsets to generate association rules

Mining Frequent Itemsets

- ⦿ **Input:** A set of transactions T , over a set of items I
- ⦿ **Output:** All itemsets with items in I having
 - ⦿ $\text{support} \geq \text{minsup}$ (support threshold)
- ⦿ Problem parameters:
 - ⦿ $N = |T|$: number of transactions
 - ⦿ $d = |I|$: number of (distinct) items
 - ⦿ w : max width of a transaction
 - ⦿ Number of possible itemsets: $M = 2^d$
- ⦿ Scale of the problem:
 - ⦿ WalMart sells 100,000 items and can store billions of baskets.
 - ⦿ The Web has billions of words and many billions of pages.

The Itemset Lattice



Representation of all possible itemsets and their relationships

Given d items, there are 2^d possible itemsets

Agenda

- ⌚ Problem definition
- ⌚ **Apriori principle, Apriori algorithm**
- ⌚ Deriving association rules from frequent itemsets
- ⌚ Reflection on Apriori algorithm

The Apriori Principle

- ⦿ Main observations: $\forall X, Y: X \subseteq Y \Rightarrow s(X) \geq s(Y)$
 - ⦿ If an itemset is frequent, so are its subsets
 - ⦿ If an itemset is infrequent, so are its supersets
- ⦿ The **Apriori** principle: **A subset of a frequent itemset must also be a frequent itemset**
 - › E.g., if $\{AB\}$ is a frequent itemset, both $\{A\}$ and $\{B\}$ must be a frequent itemset
 - ⦿ Iteratively find frequent itemsets with cardinality from 1 to m (m -itemset): Use frequent k-itemsets to explore $(k+1)$ -itemsets

Illustration of Apriori Principle

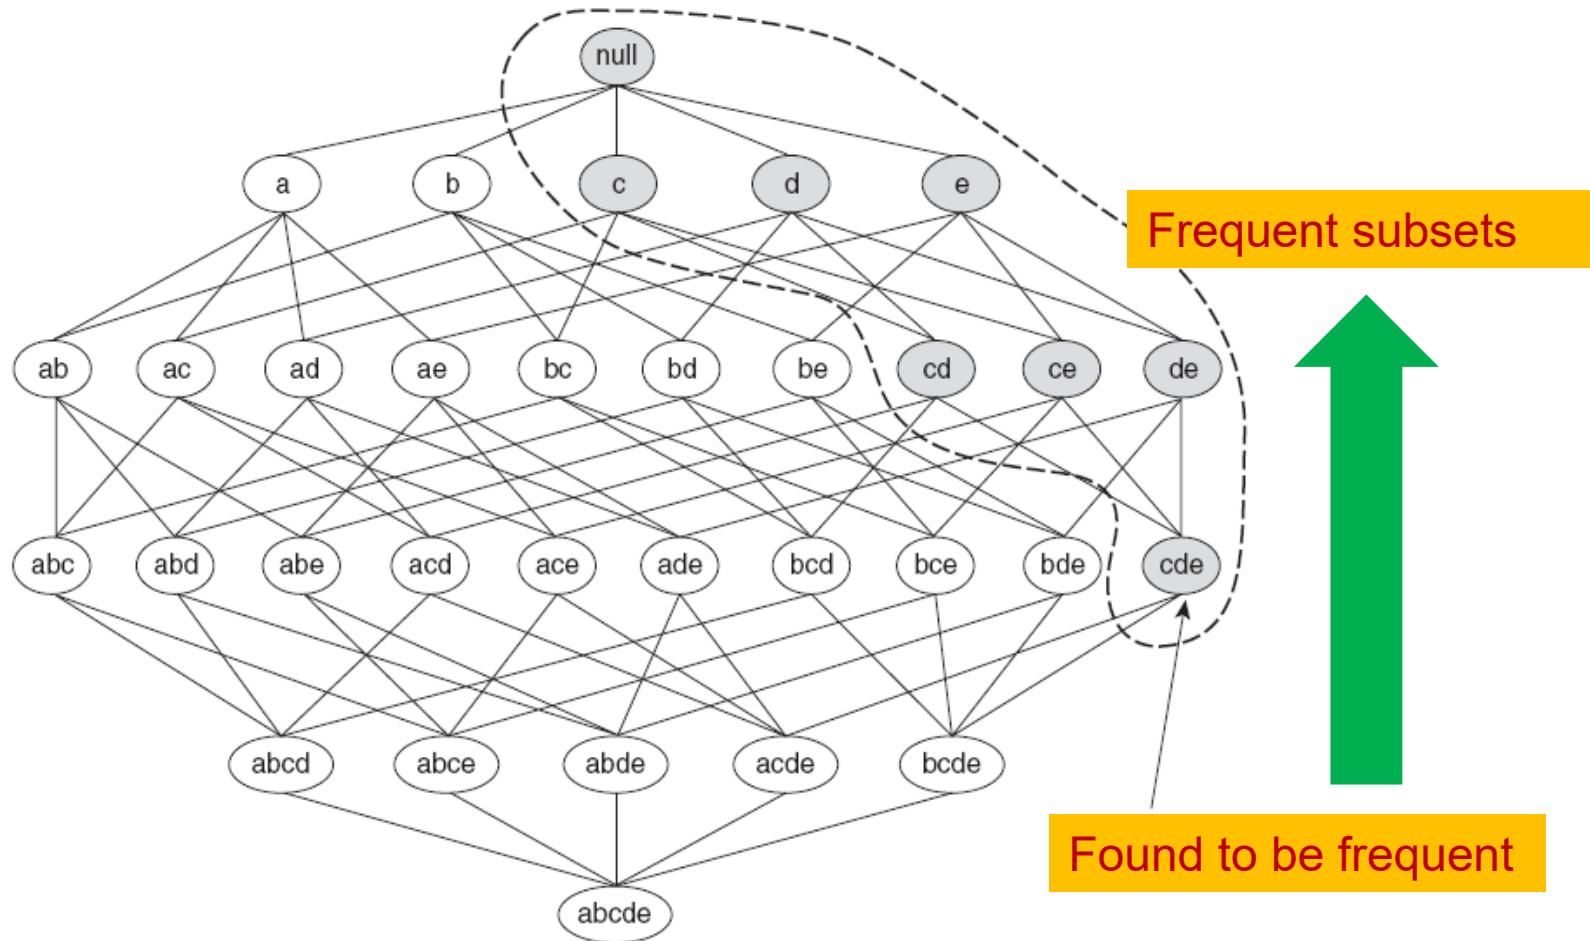
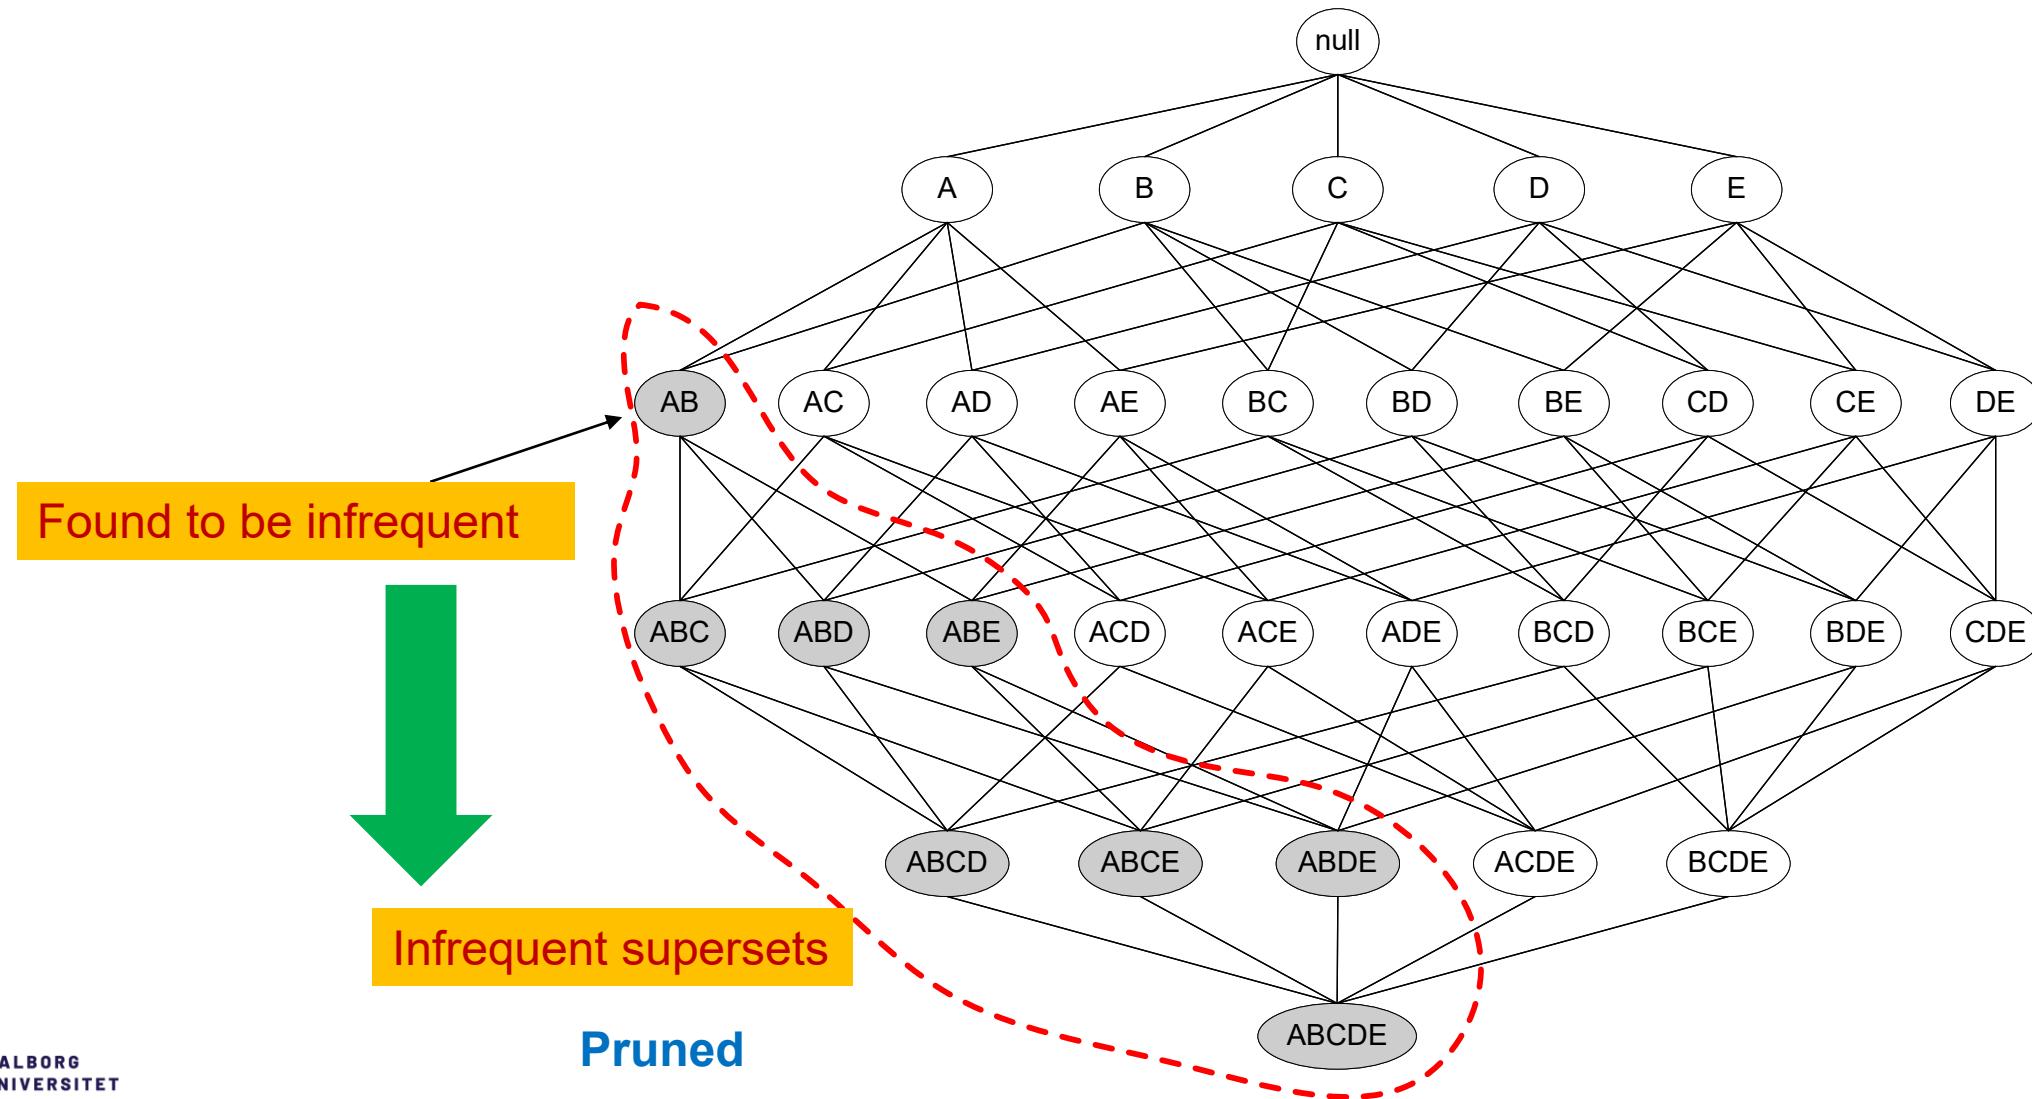


Figure 6.3. An illustration of the *Apriori* principle. If $\{c, d, e\}$ is frequent, then all subsets of this itemset are frequent.

Illustration of Apriori Principle (cont.)



Level-wise Process of Apriori Principle

Level 4 (frequent quadruples): $\{....\}$

Level 3 (frequent triplets): $\{ABD\}, \{BDF\}$

Level 2 (frequent pairs): $\{AB\}, \{AD\}, \{BD\}, \{BF\}, \{DF\}$

Level 1 (frequent items): $\{A\}, \{B\}, \{D\}, \{F\}$

Remember:

All subsets of a frequent itemset must be frequent as well

Question: Can ADF be frequent?

NO: because AF is not frequent

The Apriori Algorithm

Notations

- C_k : Candidate itemset of size k
- L_k : Frequent itemset of size k

Important steps in candidate generation

- **Prune Step:** Any k-itemset that is not frequent cannot be a subset of a frequent (k+1)-itemset
- **Join Step:** C_{k+1} is generated by joining L_k with itself

```
 $C_1 = \{\{\text{item}_1\}, \dots, \{\text{item}_N\}\};$ 
for ( $k = 1; L_k \neq \emptyset; k++$ )
    for each transaction  $t$  in transaction table  $T$ 
        increment the count of all candidates in  $C_k$  that are contained in  $t$ 
     $L_k$  = candidates in  $C_k$  with min_support (frequent)
     $C_{k+1}$  = candidates generated from  $L_k$ ;
return  $\cup_k L_k$ ;
```

Special self-join!

The Apriori Algorithm Example (1)

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

C_1
Scan T

itemset	sup.
{1}	2
{2}	3
{3}	3
{4}	1
{5}	3

L_1

itemset	sup.
{1}	2
{2}	3
{3}	3
{5}	3

min_sup=2 (or 50%)

C_2

itemset
{1 2}
{1 3}
{1 5}
{2 3}
{2 5}
{3 5}

The Apriori Algorithm Example (2)

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

$\text{min_sup}=2$ (or 50%)

L_2	itemset	sup
	{1 3}	2
	{2 3}	2
	{2 5}	3
	{3 5}	2

C_2	itemset	sup
	{1 2}	1
	{1 3}	2
	{1 5}	1
	{2 3}	2
	{2 5}	3
	{3 5}	2

← Scan T

C_2	itemset
	{1 2}
	{1 3}
	{1 5}
	{2 3}
	{2 5}
	{3 5}



C_3	itemset
	{2 3 5}

The Apriori Algorithm Example (3)

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

min_sup=2 (or 50%)



C_3

itemset
{2 3 5}

Scan T

itemset	sup
{2 3 5}	2

The Apriori Algorithm Example (4)

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

min_sup=2 (or 50%)

The result of frequent itemsets

itemset	sup.
{1}	2
{2}	3
{3}	3
{5}	3

itemset	sup
{1 3}	2
{2 3}	2
{2 5}	3
{3 5}	2

itemset	sup
{2 3 5}	2

$$L_1 \cup L_2 \cup L_3$$

Candidates Generation

- Suppose the items in L_k are listed in an order
- Step 1: self-joining L_k to get C_{k+1} (In SQL)

INSERT INTO C_{k+1}

SELECT $p.item_1, p.item_2, \dots, p.item_k, q.item_k$

FROM $L_k p, L_k q$

WHERE $p.item_1=q.item_1, \dots, p.item_{k-1}=q.item_{k-1}, p.item_k < q.item_k$

- Step 2: pruning frequent itemsets in C_{k+1}

forall ***itemsets c in C_{k+1}*** do

forall ***k-subsets s of c*** do

if (s is not in L_k) then delete c from C_{k+1}

The Previous Example

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

min_sup=2 (or 50%)

We only need
to match
 $\{2 3\}$ with $\{2 5\}$

L_2

itemset	sup
$\{1 3\}$	2
$\{2 3\}$	2
$\{2 5\}$	3
$\{3 5\}$	2

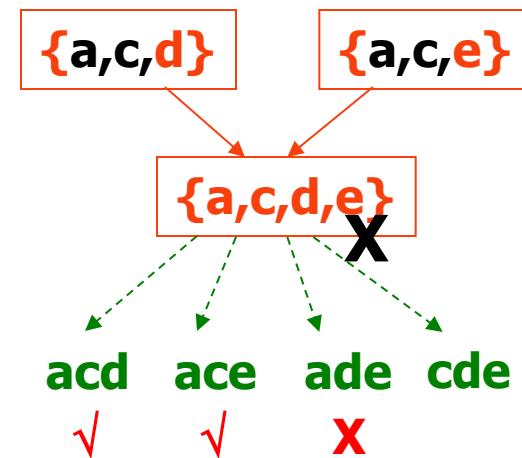


C_3

itemset
$\{2 3 5\}$

Example of Candidates Generation

- ⦿ $L_3 = \{abc, abd, acd, ace, bcd\}$
- ⦿ Self-joining: $L_3 \bowtie L_3$
 - ⦿ $abcd$ from abc and abd
 - ⦿ $acde$ from acd and ace
 - ⦿ No need to match other pairs
- ⦿ Pruning:
 - ⦿ $acde$ is removed because ade is not in L_3
- ⦿ $C_4 = \{abcd\}$
 - ⦿ Scanning transaction table T is still needed to get the frequencies for items in C_4 (to decide the correct L_4)



Agenda

- ⌚ Problem definition
- ⌚ Apriori principle, Apriori algorithm
- ⌚ Deriving association rules from frequent itemsets
- ⌚ Reflection on Apriori algorithm

Generating Association Rules from Frequent Itemsets

- Assume that we have discovered the frequent itemsets and their support
- How do we generate association rules?
- Frequent itemsets:



$l \rightarrow$

{1}	2
{2}	3
{3}	3
{5}	3
{1,3}	2
{2,3}	2
{2,5}	3
{3,5}	2
{2,3,5}	2

?

- For each frequent itemset l , find all nonempty subsets s .
- For each s , generate rule $s \Rightarrow l-s$, if $\text{sup}(l)/\text{sup}(s) \geq \text{min_conf}$

Example: $l = \{2,3,5\}$, $\text{min_conf} = 75\%$

- | | | |
|-----------------------------|--------------|---|
| $\{2,3\} \Rightarrow \{5\}$ | $2/2=100\%$ | ✓ |
| $\{2,5\} \Rightarrow \{3\}$ | $2/3=66.6\%$ | ✗ |
| $\{3,5\} \Rightarrow \{2\}$ | $2/2=100\%$ | ✓ |

do the rest as an exercise

Association Rules in Jupyter Notebook

➤ Library **mlxtend**

- To install the library: `pip install mlxtend` in Anaconda Prompt
- `from mlxtend.frequent_patterns import apriori`: frequent itemsets
- `from mlxtend.frequent_patterns import association_rules`: rules

➤ Exercises on real data

- `store_data.csv` (in Moodle)
- (7501, 20)
 - 7501 transactions, each having at most 20 items

Agenda

- ⌚ Problem definition
- ⌚ Apriori principle, Apriori algorithm
- ⌚ Deriving association rules from frequent itemsets
- ⌚ Reflection on Apriori algorithm

Performance Bottlenecks of Apriori

- Is Apriori fast enough?
- The core of the Apriori algorithm:
 - Use frequent k -itemsets to generate **candidate** frequent $(k+1)$ -itemsets
 - Use full table scan and pattern matching to collect counts for the candidate itemsets
- The bottleneck of Apriori: **candidate generation**
 - Huge candidate sets:
 - A 10^4 -sized frequent 1-itemset will generate 10^7 candidate 2-itemsets
 - To discover a frequent pattern of size 100, e.g., $\{a_1, a_2, \dots, a_{100}\}$, one needs to generate $2^{100} \approx 10^{30}$ candidates.
 - Multiple scans of database table:
 - Needs $(n + 1)$ scans, n is the length of the longest pattern

Methods to Improve Apriori's Efficiency

- ⌚ Transaction reduction
 - ⌚ A transaction that does not contain any frequent k-itemset is useless in subsequent scans
- ⌚ Partitioning
 - ⌚ Any itemset that is potentially frequent in transaction table T must be frequent in at least one of the partitions of T.

Partitioning

- ⌚ Divide transaction table T into partitions T_1, T_2, \dots, T_p such that each T_i fits in the main memory
- ⌚ Apply Apriori to each partition
- ⌚ Any frequent itemset must be frequent in at least one partition

1. Divide T into partitions T^1, T^2, \dots, T^p ;
2. For $i = 1$ to p do
3. $L^i = \text{Apriori}(T^i)$;
4. $C = L^1 \cup \dots \cup L^p$;
5. Count C on T to generate L ;

Partitioning Example (1)

	Transaction	Items
T_1	t_1	Bread,Jelly,PeanutButter
	t_2	Bread,PeanutButter
	t_3	Bread,Milk,PeanutButter
T_2	t_4	Beer,Bread
	t_5	Beer,Milk

`min_support = 10%`

$L(T_1) = \{\{Bread\}, \{Jelly\}, \{PeanutButter\}, \{Bread, Jelly\}, \{Bread, PeanutButter\}, \{Jelly, PeanutButter\}, \{Bread, Jelly, PeanutButter\}\}$

$L(T_2) = \{\{Bread\}, \{Milk\}, \{PeanutButter\}, \{Bread, Milk\}, \{Bread, PeanutButter\}, \{Milk, PeanutButter\}, \{Bread, Milk, PeanutButter\}, \{Beer\}, \{Beer, Bread\}, \{Beer, Milk\}\}$

Partitioning Example (2)

Transaction	Items
t_1	Bread,Jelly,PeanutButter
t_2	Bread,PeanutButter
t_3	Bread,Milk,PeanutButter
t_4	Beer,Bread
t_5	Beer,Milk

min_support = 10%

$$L(T_1) = \{\{Bread\}, \{Jelly\}, \\ \{PeanutButter\}, \{Bread, Jelly\}, \\ \{Bread, PeanutButter\}, \{Jelly, \\ PeanutButter\}, \\ \{Bread, Jelly, PeanutButter\}\}$$

$$L(T_2) = \{\{Bread\}, \{Milk\}, \\ \{PeanutButter\}, \{Bread, Milk\}, \\ \{Bread, PeanutButter\}, \{Milk, \\ PeanutButter\}, \\ \{Bread, Milk, PeanutButter\}, \{Beer\}, \\ \{Beer, Bread\}, \{Beer, Milk\}\}$$

$$C = L(T_1) \cup L(T_2)$$

Count itemsets in C with respect to T ,
and prune infrequent ones.

Partitioning's Pros and Cons

- ⦿ Advantages:
 - ⦿ It adapts to available main memory
 - ⦿ It can be easily parallelized
 - › Maximum number of database table scans is two (why?)
 - › One for partitioning the transaction table, and one for the final counting
- ⦿ Disadvantages:
 - ⦿ May have many candidates for the second scan
 - ⦿ A countermeasure: associate the frequency to each itemset in each partition, and the final global counting can be a simple aggregation.

More Efficient Approach: FP-tree

- Using FP-tree for finding frequent items
- Compress a large database table into a compact, Frequent-Pattern tree (FP-tree) structure
 - highly condensed, but complete for frequent pattern mining
 - avoid costly database table scans
- Develop an efficient, FP-tree-based frequent pattern mining method
 - A divide-and-conquer methodology: decompose mining tasks into smaller ones
 - Avoid candidate generation: sub-database test only!
- FP-growth: mining frequent patterns without candidate generation

Summary

- ❯ Association rule definition
 - ❯ Support, confidence, lift and association rule
 - ❯ Frequent itemsets
 - ❯ Steps for association rule mining
- ❯ Apriori algorithm
- ❯ Deriving association rules from frequent itemsets
- ❯ Criticism on Apriori

Exercises

1. Refer to the transaction table to the right. Say $\text{sup}(ab)=100$

● Determine the possible values of $\text{sup}(a)$

- Conclusion: $\text{sup}(a) \underline{\hspace{2cm}} 100$

- Hint: Is it possible that $\text{sup}(a)=70$? Why?

● Determine the possible values of $\text{sup}(abc)$

- Conclusion: $\text{sup}(abc) \underline{\hspace{2cm}} 100$

- Hint: Is it possible that $\text{sup}(abc)=120$? Why?

2. Slides 33 (Hands-on, optional)

Transaction table
(1000 rows)

TID	Items
1	a,b,c
2	a,c
3	b,e,f
...

Choose either " \leq " or " \geq "

Readings

► Mandatory reading

- ⦿ Jiawei Han, Micheline Kamber and Jian Pei. Data Mining: Concepts and Techniques (3rd edition), Elsevier Science Ltd, 2011.
 - › Chapter 6

► Further readings

- ⦿ Rakesh Agrawal, Ramakrishnan Srikant: Fast Algorithms for Mining Association Rules in Large Databases. VLDB 1994: 487-499
- ⦿ Jiawei Han, Jian Pei, Yiwen Yin: Mining Frequent Patterns without Candidate Generation. SIGMOD 2000: 1-12

► Acknowledgment: Slides are from

- ⦿ Margaret H. Dunham (Data Mining: Introductory and Advanced Topics, Prentice Hall, 2002)
- ⦿ The HKP textbook
- ⦿ Man Lung Yiu and Panagiotis Karras

Readings for Coding

⦿ Mandatory readings

- ⦿ Association Rule: <https://www.geeksforgeeks.org/association-rule/?ref=lbp>
- ⦿ Frequent Itemsets: <https://www.geeksforgeeks.org/frequent-item-set-in-data-set-association-rule-mining/?ref=lbp>
- ⦿ Apriori Algorithm: <https://www.geeksforgeeks.org/apriori-algorithm/?ref=lbp>

⦿ Further readings

- ⦿ Documentation of mlxtend's frequent
 - › http://rasbt.github.io/mlxtend/user_guide/frequent_patterns/apriori/
 - › http://rasbt.github.io/mlxtend/user_guide/frequent_patterns/association_rules/
- ⦿ Tutorials
 - › <https://www.geeksforgeeks.org/implementing-apriori-algorithm-in-python/>
 - › <https://www.kaggle.com/code/annettecatherinepaul/apriori-algorithm-association-rule-mining>
 - › <https://towardsdatascience.com/understand-and-build-fp-growth-algorithm-in-python-d8b989bab342> (FP-Growth, advanced)