



# Data Intensive Systems (DIS)

## KBH-SW7 E25

### 3. Distributed Databases



AALBORG  
UNIVERSITY

# Agenda

- ⌚ Introduction
  - ⌚ Heterogeneous and Homogeneous Databases
- ⌚ Distributed Data Storage
- ⌚ Distributed Query Processing and Optimization
- ⌚ Distributed Transactions

# Distributed Database System

- A distributed database system consists of loosely coupled sites that **share no physical component**.
- Database systems that run on different *sites* are independent of each other.
- Transactions and queries may access data at one or more sites.

# Parallel DB vs. Distributed DB

- ▶ Parallel Databases
  - ▶ Machines are physically close to each other, e.g., same server room
  - ▶ Machines connect with dedicated high-speed LANs and switches, and thus communication cost is assumed to be small
  - ▶ Can share memory, share disk, or share nothing
- ▶ Distributed Databases
  - ▶ Machines can be far away from each other, e.g., in different continents
  - ▶ Can be connected using public-purpose network, e.g., Internet
  - ▶ Communication cost and problems cannot be ignored
  - ▶ Usually built using shared-nothing architecture

# Distributed Databases

- In a **homogeneous** distributed database
  - All sites have identical software
  - All sites are aware of each other and agree to cooperate in processing user requests.
  - Each site surrenders part of its autonomy in terms of right to change schemas or software.
  - Appears to user as a single system.
- In a **heterogeneous** distributed database
  - Different sites may use different schemas and software.
    - Difference in schema is a major problem for query processing.
      - Data integration, schema matching
    - Difference in software is a major problem for transaction processing.
  - Sites may not be aware of each other and may provide only limited facilities for cooperation in transaction processing.

# Agenda

- › Introduction
- › **Distributed Data Storage**
- › Distributed Query Processing and Optimization
- › Distributed Transactions

# Distributed Data Storage

- Distributed Database Design:
  - How should the database and applications on top of it be placed *across the sites*?
  - We assume relational data model, but the techniques can be applicable to others.
- Replication
  - System maintains multiple copies of data, stored in different sites, for **faster retrieval and fault tolerance**.
- Fragmentation (partitioning)
  - A relation is partitioned into several fragments stored in distinct sites.
- Replication and fragmentation can be combined.
  - Relation is partitioned into several fragments: system maintains several identical replicas of each such fragment.

# Data Replication

- A relation or fragment of a relation is **replicated** if it is stored redundantly in two or more sites.
- **Full replication** of a relation is the case where the relation is stored at all sites.
- **Fully redundant** databases are those in which every site contains a copy of the entire database.
- Advantages of Replication
  - **Availability**: failure of site containing relation  $r$  does not result in unavailability of  $r$  if replicas exist.
  - **Parallelism**: queries on  $r$  may be processed by several sites in parallel.
  - **Reduced data transfer**: relation  $r$  is available locally at each site that has a replica of  $r$ .

# Data Replication, cont.

## ➤ Disadvantages of Replication

### ● Increased cost of updates

- Each replica of relation  $r$  must be updated.

### ● Increased complexity of concurrency control

- Concurrent updates to distinct replicas may lead to inconsistent data unless special concurrency control mechanisms are implemented.

- E.g., consider money transfer between replicates on different sites.

- One solution: choose one copy as **primary copy** and apply concurrency control operations on primary copy

# Data Fragmentation

- ⌚ A relation  $r$  is partitioned into fragments  $r_1, r_2, \dots, r_n$  which contain sufficient information to reconstruct relation  $r$ .
- ⌚ **Horizontal fragmentation**: each tuple of  $r$  is assigned to one or more fragments
- ⌚ **Vertical fragmentation**: the schema for relation  $r$  is split into several smaller schemas
  - ⌚ All schemas must contain a common key to ensure *lossless join property*.
  - ⌚ A special attribute, e.g., the tuple-id attribute, may be added to each schema to serve as the key.

# Example Relations

## ➤ Relation *account*

branch_name	account_number	balance
Brussels	123456789	10000

## ➤ Relation *employee*

account_number	balance	branch_name	customer_name
123456789	10000	Brussels	John Doe

# Horizontal Fragmentation of *account* Relation

<b>branch_name</b>	<b>account_number</b>	<b>balance</b>
Hillside	A-305	500
Hillside	A-226	336
Hillside	A-155	62

$account_1 = \sigma_{branch\_name="Hillside"}(account)$

<b>branch_name</b>	<b>account_number</b>	<b>balance</b>
Valleyview	A-177	205
Valleyview	A-402	10000
Valleyview	A-408	1123
Valleyview	A-639	750

$account_2 = \sigma_{branch\_name="Valleyview"}(account)$

# Vertical Fragmentation of *employee* Relation

branch_name	customer_name	tuple_id
Hillside	Lowman	1
Hillside	Camp	2
Valleyview	Camp	3
Valleyview	Kahn	4
Hillside	Kahn	5
Valleyview	Kahn	6
Valleyview	Green	7

$$deposit_1 = \Pi_{branch\_name, customer\_name, tuple\_id} (\text{employee})$$

account_number	balance	tuple_id
A-305	500	1
A-226	336	2
A-177	205	3
A-402	10000	4
A-155	62	5
A-408	1123	6
A-639	750	7

$$deposit_2 = \Pi_{account\_number, balance, tuple\_id} (\text{employee})$$

# Advantages of Fragmentation

- ⦿ Horizontal fragmentation:
  - ⦿ allows parallel processing on fragments of a relation
  - ⦿ allows a relation to be split so that tuples are located where they are most frequently accessed
- ⦿ Vertical fragmentation:
  - ⦿ allows parallel processing on a relation
  - ⦿ allows tuples to be split so that each part of the tuple is stored where it is most frequently accessed
  - ⦿ tuple-id attribute allows efficient joining of vertical fragments
- ⦿ Vertical and horizontal fragmentation can be mixed.
  - ⦿ Fragments may be successively fragmented to an arbitrary depth.

# Data Transparency

- Data transparency
  - Degree to which a system user may remain *unaware* of the details of *how* and *where* the data items are stored in a distributed system.
- Consider transparency issues in relation to:
  - Fragmentation transparency
  - Replication transparency
  - Location transparency
- NB: Similar ideas are used in distributed file systems like HDFS

# Agenda

- ⌚ Introduction
- ⌚ Distributed Data Storage
- ⌚ **Distributed Query Processing and Optimization**
  - ⌚ Fragmented storage
  - ⌚ Non-fragmented storage
- ⌚ Distributed Transactions

# Distributed Query Processing

- For centralized systems, the primary criterion for measuring the cost of a particular query strategy is the number of disk accesses (IOs).
- In a distributed system, other issues must be considered:
  - Data transmission cost over the network (often critical).
  - The potential gain in performance: several sites process parts of the query in parallel.
- Translating algebraic queries on fragments
  - It must be possible to construct relation  $r$  from its fragments.
  - Replace relation  $r$  by the expression to construct relation  $r$  from its fragments.

# Query Transformation Example

- Consider the horizontal fragmentation of the *account* relation

$\text{account}_1 = \sigma_{\text{branch\_name} = \text{"Hillside"}}(\text{account})$

$\text{account}_2 = \sigma_{\text{branch\_name} = \text{"Valleyview"}}(\text{account})$

- The query  $\sigma_{\text{branch\_name} = \text{"Hillside"}}(\text{account})$  becomes

$\sigma_{\text{branch\_name} = \text{"Hillside}}(\text{account}_1 \cup \text{account}_2)$

which is optimized into

$\sigma_{\text{branch\_name} = \text{"Hillside}}(\text{account}_1) \cup \sigma_{\text{branch\_name} = \text{"Hillside}}(\text{account}_2)$

- Since  $\text{account}_1$  has only tuples of the Hillside branch, we can eliminate the selection operation.

- Apply the definition of  $\text{account}_2$  to obtain

$\sigma_{\text{branch\_name} = \text{"Hillside}}(\sigma_{\text{branch\_name} = \text{"Valleyview}}}(\text{account}))$

- This expression always gives an empty set.

- Finally, the Hillside site returns its whole  $\text{account}_1$  as the result of the query.

# Reduced Query due to Fragmentation

Two relations with corresponding schemas:

**E(Eno, Ename, Title)**      **G(Eno, Jno, Resp, Dur)**

- E is *horizontally* fragmented into  $E_1$ ,  $E_2$  and  $E_3$  as follows:
  - ⦿  $E_1 = \sigma_{Eno \leq e3}(E)$
  - ⦿  $E_2 = \sigma_{e3 < Eno \leq e6}(E)$
  - ⦿  $E_3 = \sigma_{Eno > e6}(E)$
- G is *horizontally* fragmented into  $G_1$  and  $G_2$  as follows:
  - ⦿  $G_1 = \sigma_{Eno \leq e3}(G)$
  - ⦿  $G_2 = \sigma_{Eno > e3}(G)$
- Consider the following *natural join* query-1:

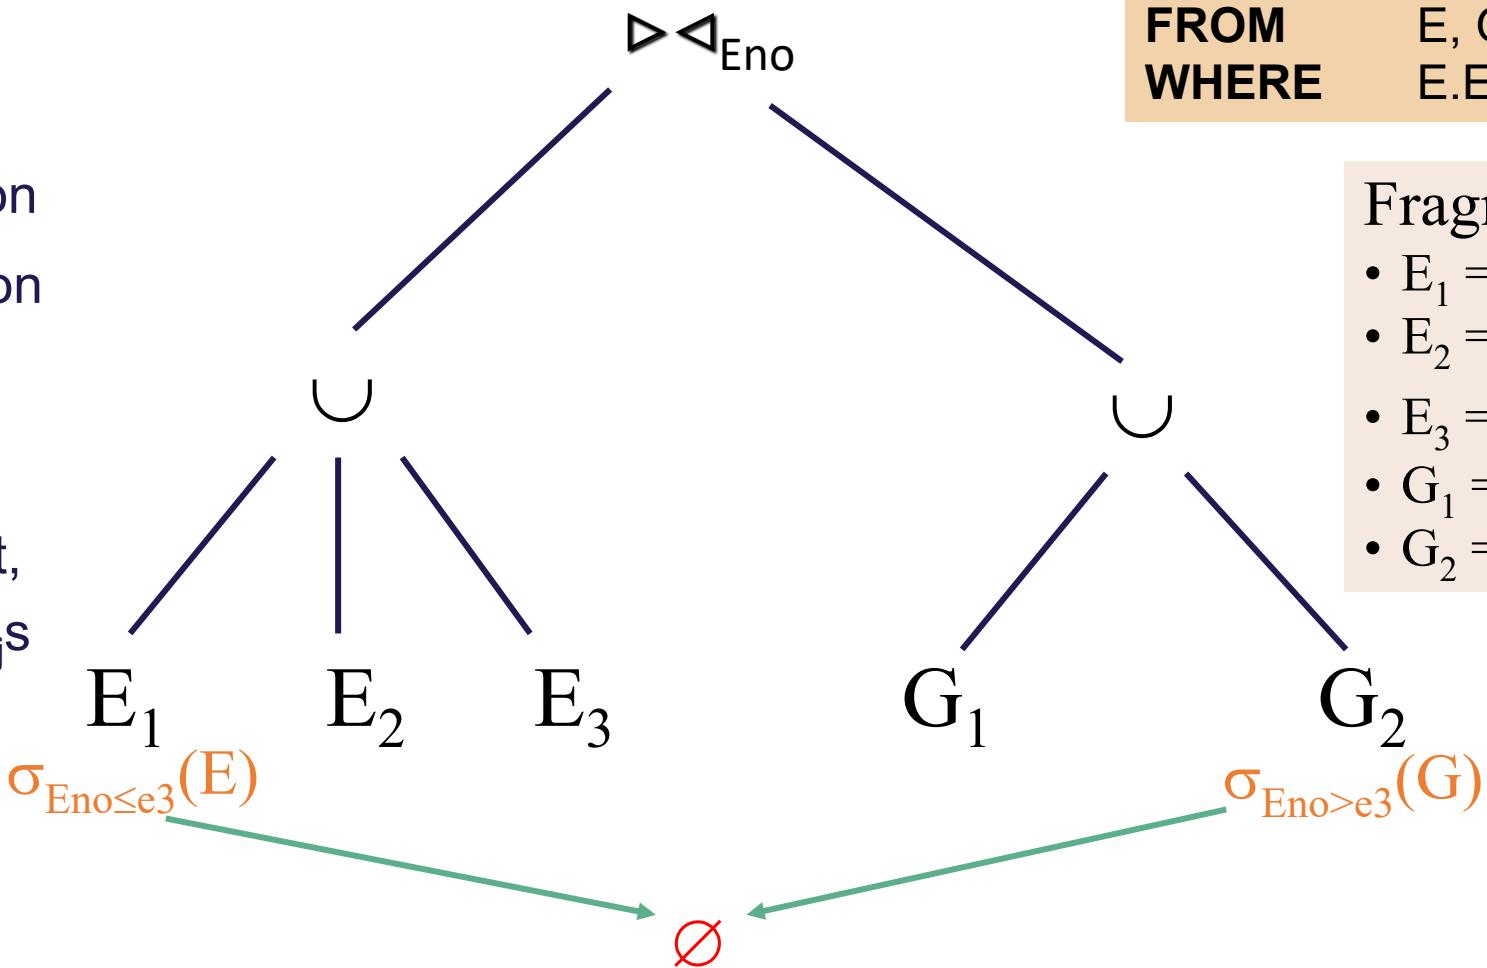
```
SELECT *
FROM E, G
WHERE E.Eno = G.Eno
```

# Query Tree for query-1:

## Steps

1. Get E by union
2. Get G by union
3. Join E and G

Steps 1 and 2 can be parallelized. But, not all  $E_i$ s match  $G_j$ s



Empty joins

- $E_1 \bowtie G_2$
- $E_2 \bowtie G_1$
- $E_3 \bowtie G_1$

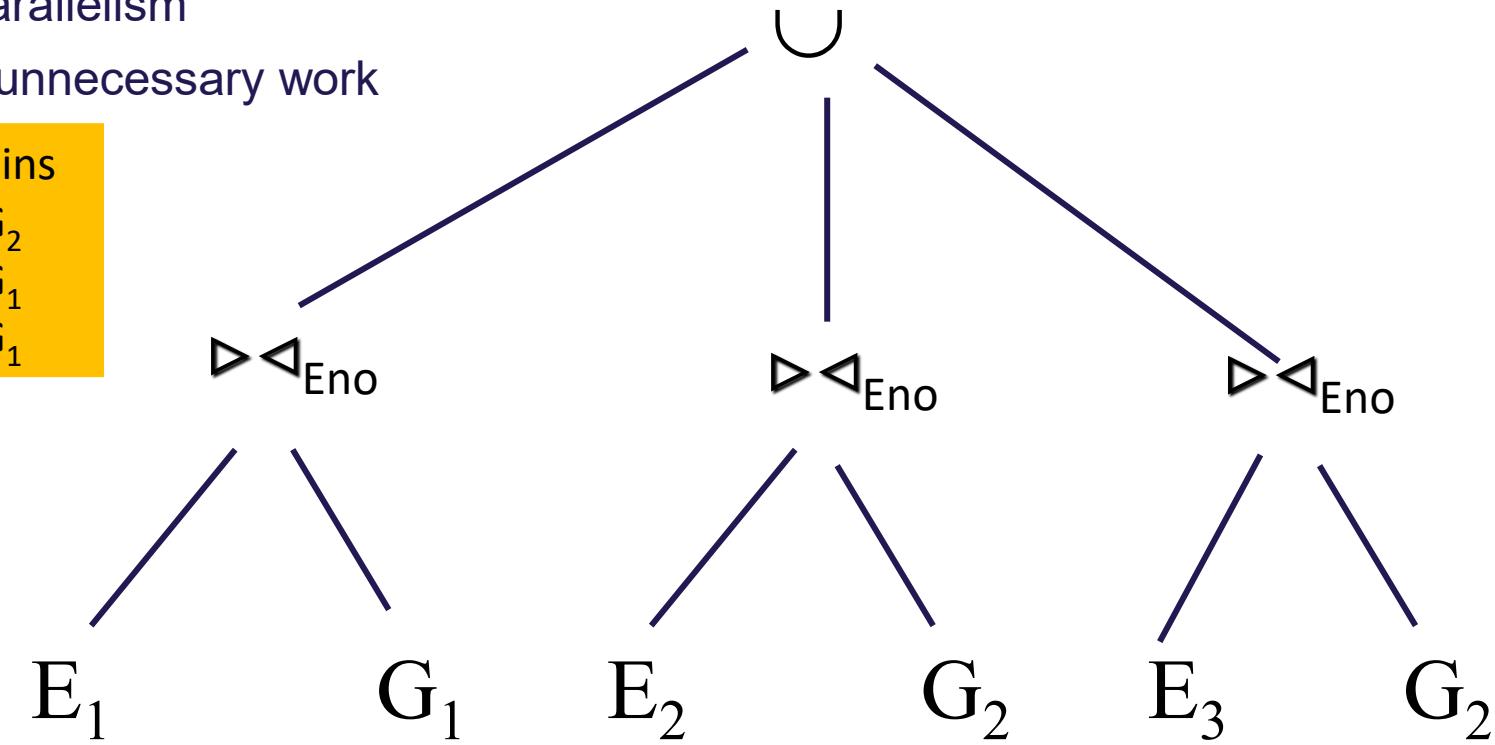
# Reduced Query Tree for query-1:

- Advantages of the reduced query tree:

- Provide parallelism
- Eliminate unnecessary work

Empty joins

- $E_1 \bowtie G_2$
- $E_2 \bowtie G_1$
- $E_3 \bowtie G_1$



```
SELECT *  
FROM E, G  
WHERE E.Eno = G.Eno
```

Fragments

- $E_1 = \sigma_{Eno \leq e3}(E)$
- $E_2 = \sigma_{e3 < Eno \leq e6}(E)$
- $E_3 = \sigma_{Eno > e6}(E)$
- $G_1 = \sigma_{Eno \leq e3}(G)$
- $G_2 = \sigma_{Eno > e3}(G)$

# Another Reduced Query Example

**E(Eno, Ename, Title)**

**G(Eno, Jno, Resp, Dur)**

Fragmentation is as follows:

$$E_1 = \sigma_{Title='Programmer'}(E)$$

$$E_2 = \sigma_{Title \neq 'Programmer'}(E)$$

$G_1 = G \bowtie_{Eno} E_1$ : Those G records that contribute to  $G \bowtie E_1$ .

$G_2 = G \bowtie_{Eno} E_2$ : Those G records that contribute to  $G \bowtie E_2$ .

Query-2:      **SELECT**                   \*

**FROM**                         E, G

**WHERE**                        E.Eno = G.Eno AND

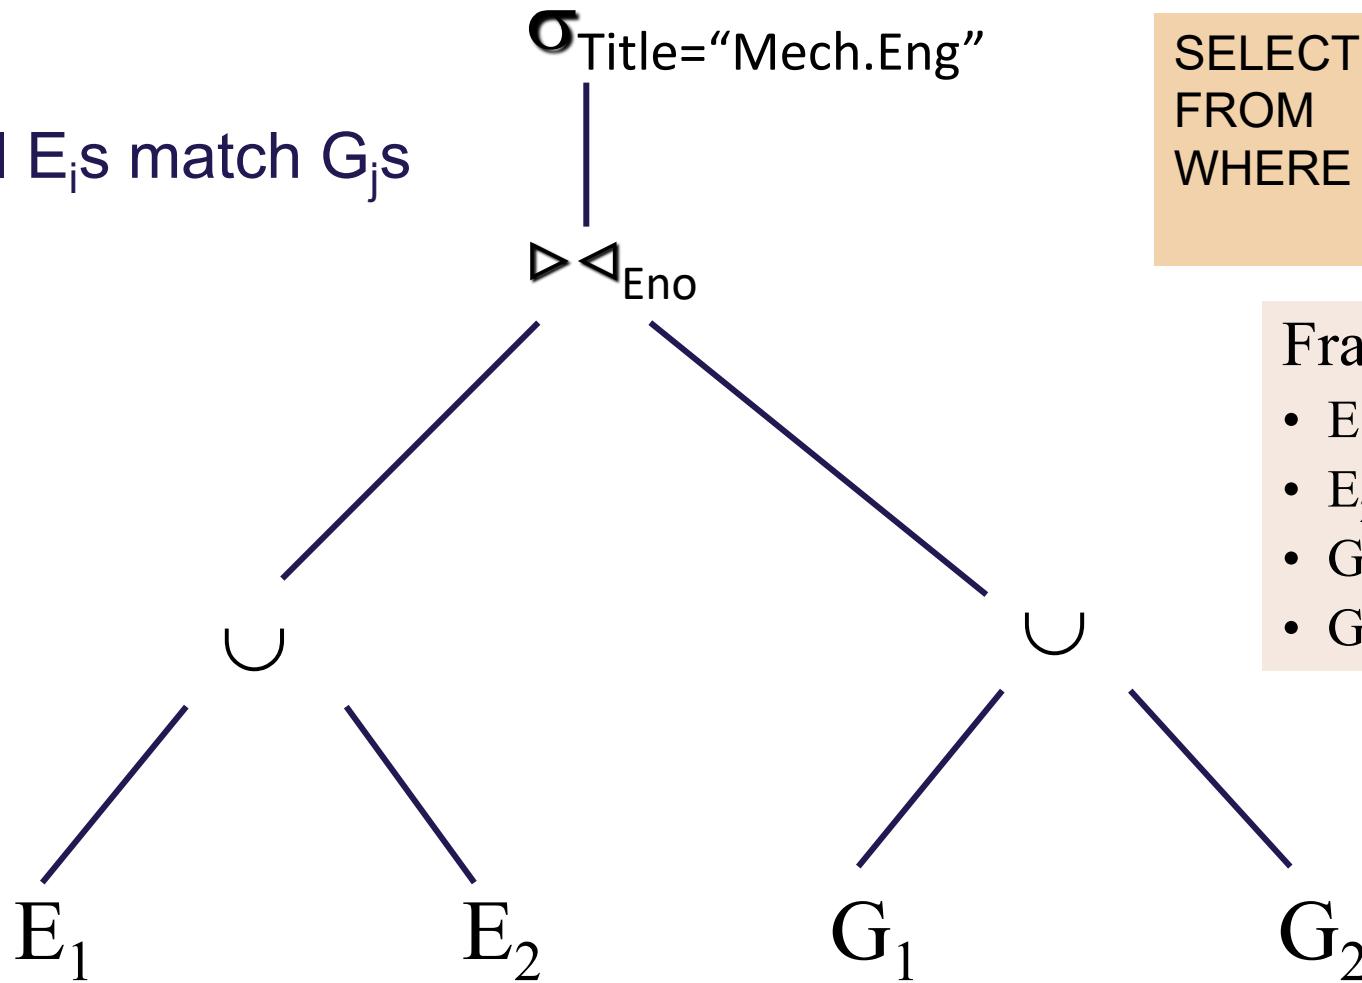
    E.Title = "Mech. Eng."

## Left semijoin $R \bowtie S$

- The result is the set of all records in R for which there is a record in S that is equal on their common attributes.
- Unlike a natural join, the other columns of S do not appear in the result.

# Query Tree for query-2:

- Again, not all  $E_i$ s match  $G_j$ s



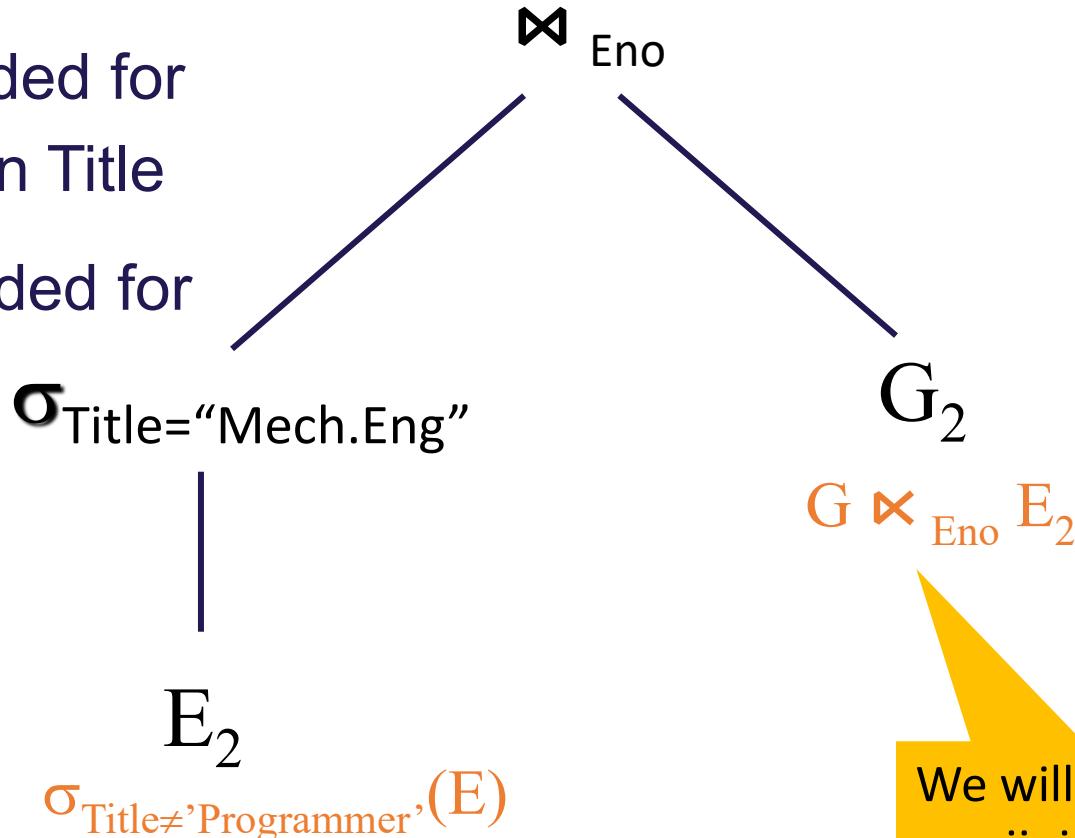
```
SELECT *  
FROM E, G  
WHERE E.Eno = G.Eno AND  
E.Title = "Mech. Eng."
```

## Fragments

- $E_1 = \sigma_{Title='Programmer'}(E)$
- $E_2 = \sigma_{Title \neq 'Programmer'}(E)$
- $G_1 = G \bowtie_{Eno} E_1$
- $G_2 = G \bowtie_{Eno} E_2$

# Reduced Query Tree for query-2:

- Only  $E_2$  is needed for the selection on Title
- Only  $G_2$  is needed for the join



```
SELECT      *
FROM        E, G
WHERE      E.Eno = G.Eno AND
          E.Title = "Mech. Eng."
```

## Fragments

- $E_1 = \sigma_{Title='Programmer},(E)$
- $E_2 = \sigma_{Title \neq 'Programmer},(E)$
- $G_1 = G \bowtie_{Eno} E_1$
- $G_2 = G \bowtie_{Eno} E_2$

# Distributed Join Processing

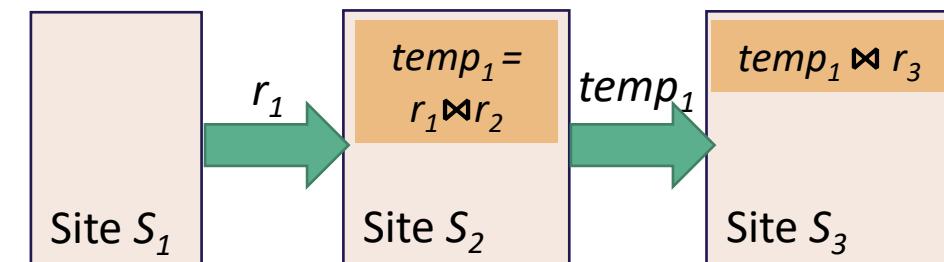
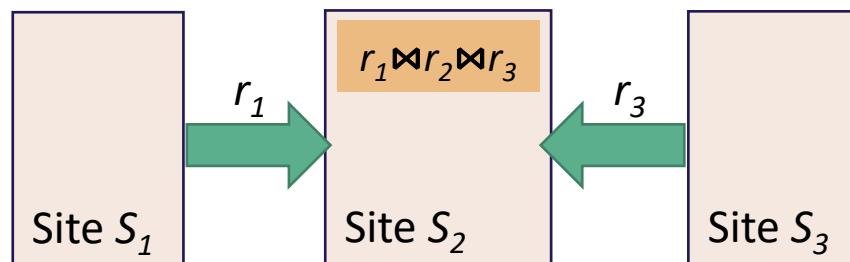
- ➊ Distributed query processing must consider following factors:
  - ➊ amount of data to be shipped between sites
  - ➋ cost of transmitting a data block between sites
  - ➌ relative processing speed at each site
- ➋ Consider the following relational algebra expression in which the three relations are *neither replicated nor* fragmented:

*account*  $\bowtie$  *depositor*  $\bowtie$  *branch*

  - ➊ *account* is stored at site  $S_1$
  - ➋ *depositor* at  $S_2$
  - ➌ *branch* at  $S_3$
- ➌ For a query issued at site  $S_i$ , the system needs to produce the result at site  $S_i$

# Simple Join Processing Strategies

- Ship copies of all three relations to site  $S_1$  and choose a strategy for processing the entire locally at site  $S_1$ . (centralized)
- Ship a copy of the account relation to site  $S_2$  and compute  $temp_1 = account \bowtie depositor$  at  $S_2$ . Ship  $temp_1$  from  $S_2$  to  $S_3$ , and compute  $temp_2 = temp_1 \bowtie branch$  at  $S_3$ . Ship the result  $temp_2$  to  $S_1$ . (serialized)
- Similar strategies, exchanging the roles  $S_1, S_2, S_3$ .

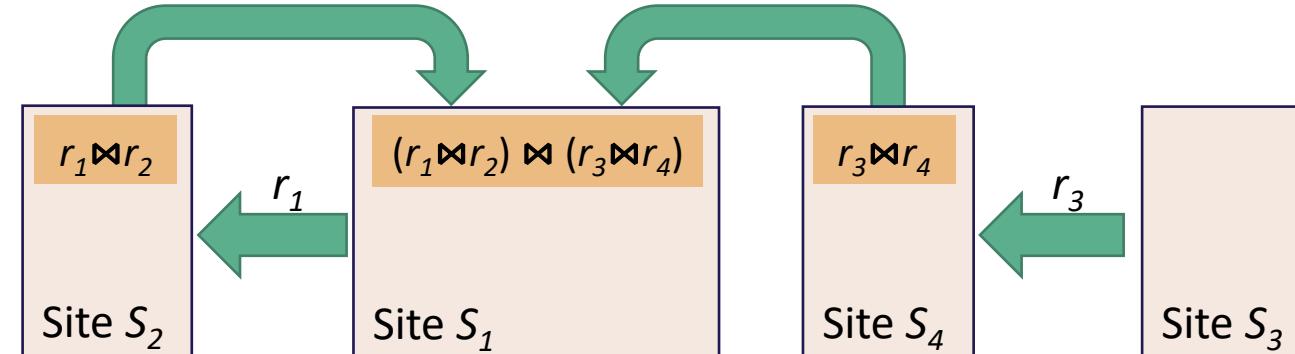


# Parallel Distributed Join Strategy

Consider  $r_1 \bowtie r_2 \bowtie r_3 \bowtie r_4$  where each relation  $r_i$  is stored at site  $S_i$ . Suppose the result must be presented at site  $S_1$ .

- ➊  $r_1$  is shipped to  $S_2$  and  $r_1 \bowtie r_2$  is computed at  $S_2$ ; simultaneously,  $r_3$  is shipped to  $S_4$  and  $r_3 \bowtie r_4$  is computed at  $S_4$ .
- ➋  $S_2$  sends tuples of  $(r_1 \bowtie r_2)$  to  $S_1$  as *they are being produced*;  $S_4$  sends tuples of  $(r_3 \bowtie r_4)$  to  $S_1$  as *they are being produced*.
- ➌ Once tuples of  $(r_1 \bowtie r_2)$  and  $(r_3 \bowtie r_4)$  arrive at  $S_1$ ,  $(r_1 \bowtie r_2) \bowtie (r_3 \bowtie r_4)$  is computed *in parallel* with the computation of  $(r_1 \bowtie r_2)$  at  $S_2$  and the computation of  $(r_3 \bowtie r_4)$  at  $S_4$ .

- ➍ Pipeline
- ➎ Parallelism

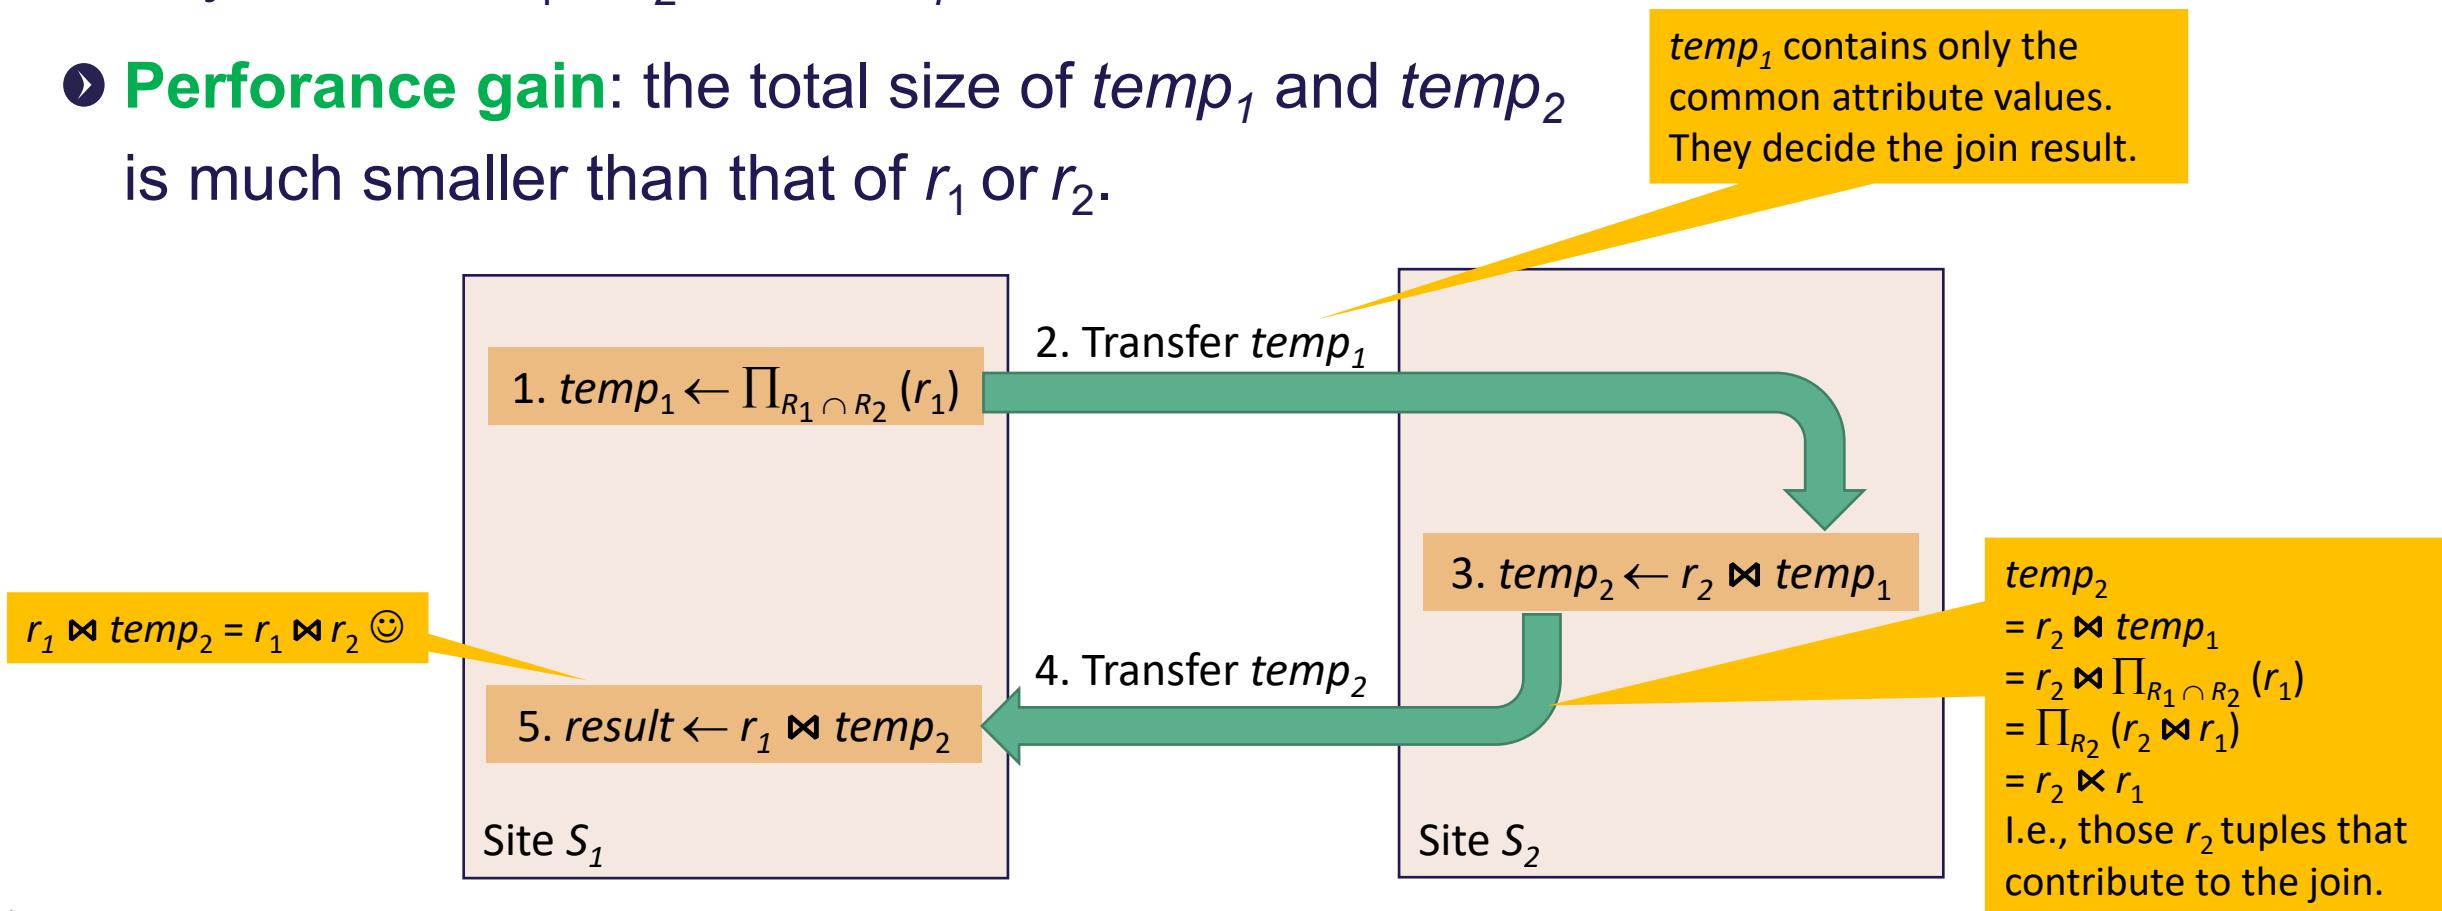


# Semijoin Strategy

- Let  $r_1$  be a relation with schema  $R_1$  stores at site  $S_1$ .  
    Let  $r_2$  be a relation with schema  $R_2$  stores at site  $S_2$ .
- Evaluate the expression  $r_1 \bowtie r_2$  and obtain the result at  $S_1$ .
- **Motivation:** Shipping whole  $r_1$  or  $r_2$  can be very expensive!
  1. Compute  $temp_1 \leftarrow \Pi_{R_1 \cap R_2}(r_1)$  at  $S_1$ .
  2. Ship  $temp_1$  from  $S_1$  to  $S_2$ .
  3. Compute  $temp_2 \leftarrow r_2 \bowtie temp_1$  at  $S_2$
  4. Ship  $temp_2$  from  $S_2$  to  $S_1$ .
  5. Compute  $r_1 \bowtie temp_2$  at  $S_1$ . This is the same as  $r_1 \bowtie r_2$ .

# Illustration of Semijoin Strategy

- Say we need  $r_1 \bowtie r_2$  at site  $S_1$ .
- Performance gain:** the total size of  $\text{temp}_1$  and  $\text{temp}_2$  is much smaller than that of  $r_1$  or  $r_2$ .



# Formal Definition of Semijoin

- ⦿ The **semijoin** of  $r_i$  with  $r_j$  is denoted by:
  - ⦿  $r_i \ltimes r_j = \Pi_{R_i} (r_i \bowtie r_j)$ , where  $R_i$  is  $r_i$ 's schema.
  - ⦿ Thus,  $r_i \ltimes r_j$  selects only those tuples of  $r_i$  that contribute to  $r_i \bowtie r_j$ .
- ⦿ In step 3 on the previous slide
  - ⦿  $\text{temp}_2 = r_2 \bowtie \text{temp}_1 = r_2 \bowtie \Pi_{R_1 \cap R_2} (r_1) = \Pi_{R_2} (r_2 \bowtie r_1) = r_2 \ltimes r_1$
  - ⦿ That's why this strategy gets its name
- ⦿ If we need  $r_1 \ltimes r_2$  at site  $S_2$ , then we need to initiate the whole process at  $S_2$  s.t. it will receive  $r_1 \ltimes r_2$ .
- ⦿ For joins of several relations, the strategy can be extended to a series of semijoin steps.

# Distributed Join Example: Naïve Strategies

- Consider the following setting and statistics:
- $r_1$  with schema R(A, B) at site  $S_1$   
 $\text{card}(R) = 10,000$
- $r_2$  with schema S(B, C) at site  $S_2$   
 $\text{card}(S) = 50,000$
- Attribute data value size:
  - A: 116 bytes; B: 4 bytes; C: 76 bytes.
- Basic costs:
  - $c_0$ : initial set-up cost between any two sites = 10
  - $c_1$ : transmission cost per data unit between any two sites = 1 per 1000 bytes
- Two naïve strategies and their costs
  - Join at  $S_1$ : Cost =  $10 + (76+4)*50000/1000 = 4010$  (shipping  $r_2$  to  $S_1$ )
  - Join at  $S_2$ : Cost =  $10 + (116+4)*10000/1000 = 1210$  (shipping  $r_1$  to  $S_2$ )

Shipping the  
smaller relation ☺

# Distributed Join Example: Semijoin Strategy

Consider the following setting and statistics:

- $r_1$  with schema R(A, B) at site  $S_1$ ,  
 $\text{card}(R) = 10,000$
- ① •  $\text{card}(\Pi_B(r_1)) = 2,000$ 
  - #of distinct values on the join attribute
- ② •  $\text{card}(r_1 \ltimes r_2) = 2,500$ 
  - #of  $r_1$  tuples that contribute to the join
- $r_2$  with schema S(B, C) at site  $S_2$   
 $\text{card}(S) = 50,000$
- ① •  $\text{card}(\Pi_B(r_2)) = 5,000$ 
  - #of distinct values on the join attribute
- ② •  $\text{card}(r_2 \ltimes r_1) = 5,000$ 
  - #of  $r_2$  tuples that contribute to the join
- Attribute data value size:
  - A: 116 bytes; B: 4 bytes; C: 76 bytes.
- Basic costs:
  - $c_0$ : initial set-up cost between any two sites = 10
  - $c_1$ : transmission cost per data unit between any two sites = 1 per 1000 bytes
- Semijoin strategy costs
  - **Result at  $S_1$ :** Cost =  $10 + 4*2000/1000 + 10 + (76+4)*5000/1000 = 428$
  - **Result at  $S_2$ :** Cost =  $10 + 4*5000/1000 + 10 + (116+4)*2500/1000 = 340$

**Semijoin cardinality**

# Agenda

- ⌚ Introduction
- ⌚ Distributed Data Storage
- ⌚ Distributed Query Processing and Optimization
- ⌚ **Distributed Transactions**

# Distributed Transactions

- Transaction may access data at several sites.
- Each site has a local **transaction manager** that is responsible for:
  - ⌚ *Maintaining* a log for recovery purposes.
  - ⌚ *Participating* in coordinating the concurrent execution of the transactions executing *at that site*.
- Each site has a **transaction coordinator** that is responsible for:
  - ⌚ *Starting* the execution of a transaction that originates at the site.
  - ⌚ *Breaking* the transaction to a number of subtransactions.
  - ⌚ *Distributing* the subtransactions to the appropriate sites for execution.
  - ⌚ *Coordinating* the termination of the transaction, which may result in the transaction being committed at all sites or aborted at all sites.

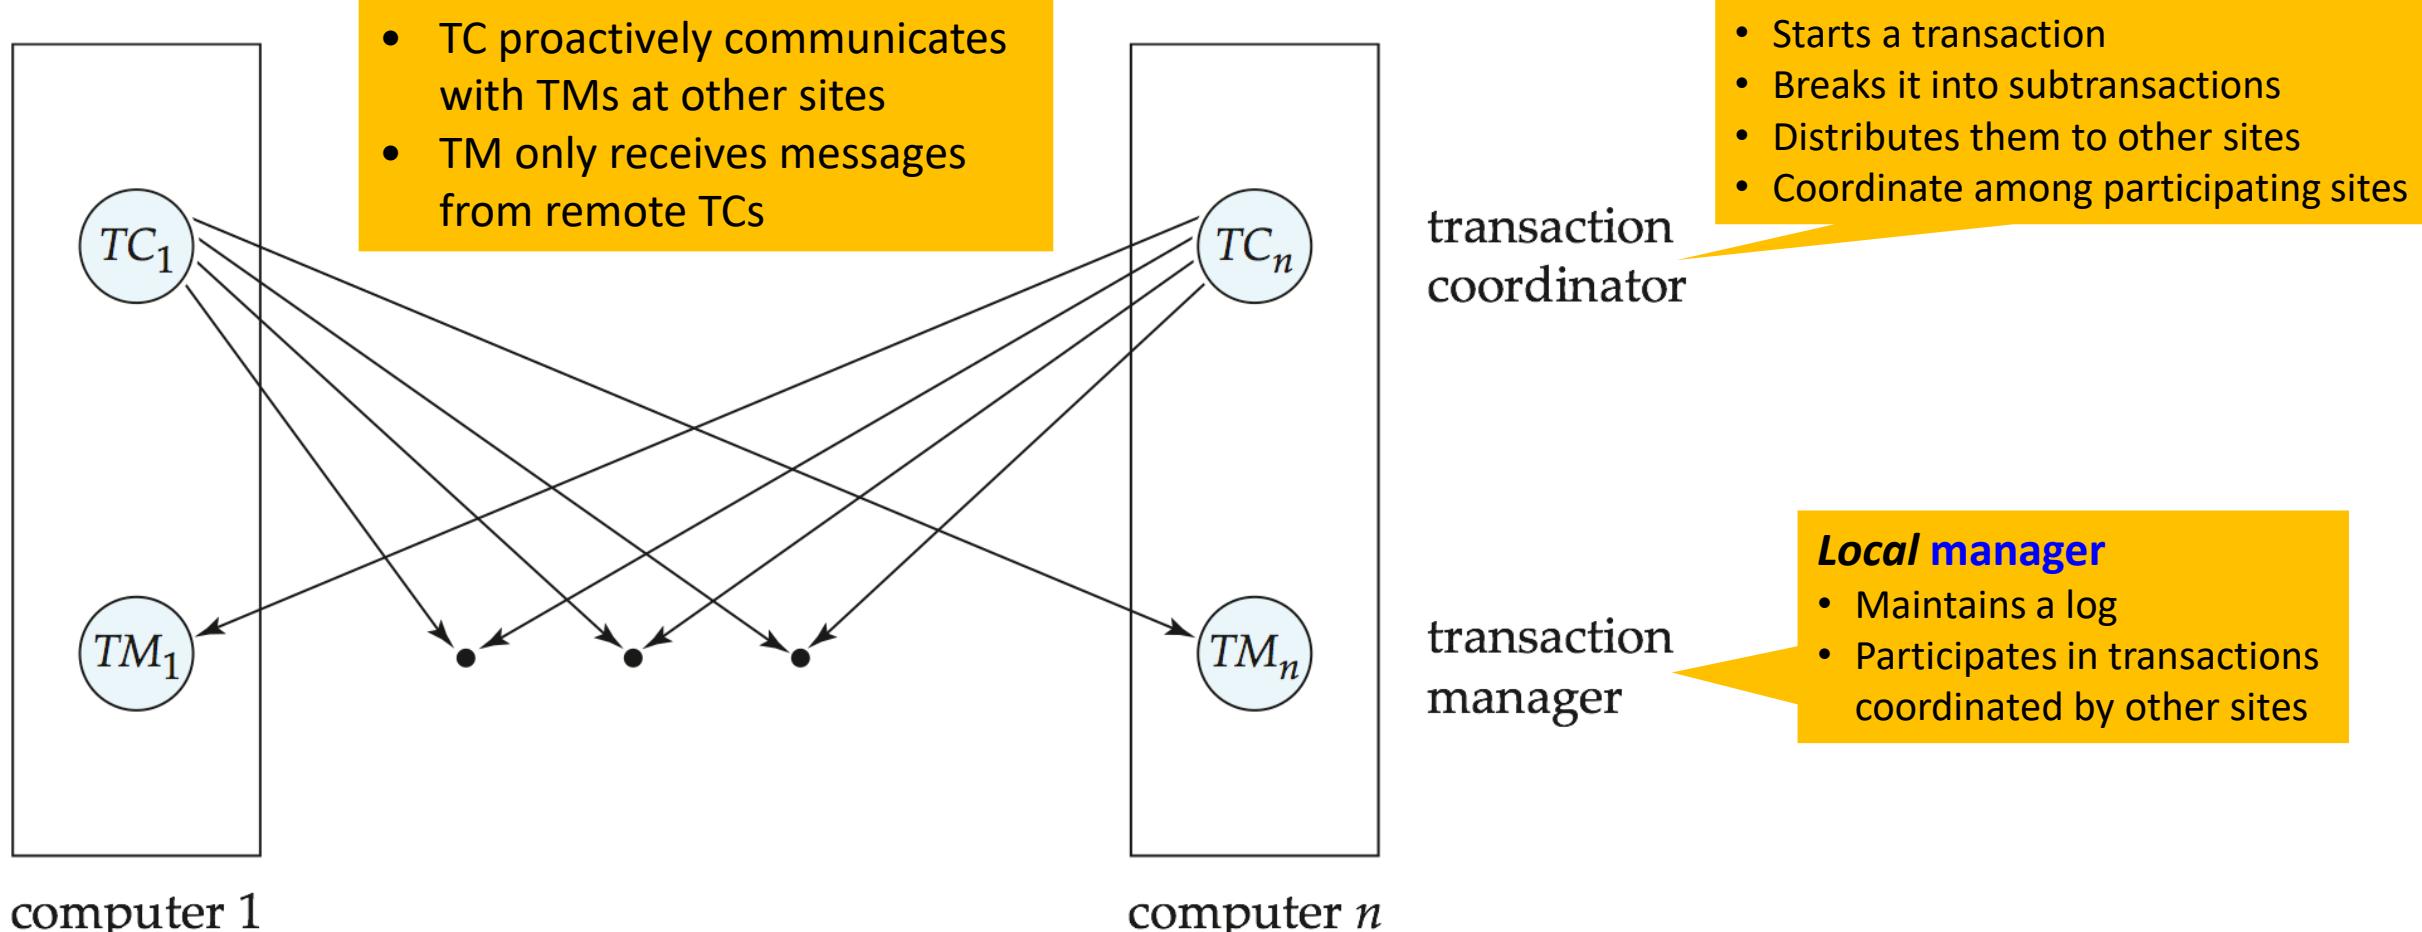
Transaction Manager

- Local role
- Used in centralized DBS

Transaction Coordinator

- Global role
- Not used in centralized DBS

# Distributed Transaction System Architecture



# System Failures

- ➊ Failures unique to distributed systems:
  - ➊ Failure of a site.
  - ➋ Loss of messages
    - › Handled by network transmission control protocols such as TCP/IP
  - ➌ Failure of a communication link
    - › Handled by network protocols, by routing messages via alternative links
  - ➍ Network partition
    - › A network is said to be **partitioned** when it has been split into two or more subsystems that lack any connection between them
      - › Note: a subsystem may consist of a single node
- ➋ Network partitioning and site failures are generally *indistinguishable*.

# Commit Protocols

- Commit protocols are used to ensure *atomicity* across sites
  - A transaction which executes at multiple sites must either be *committed at all* the sites or *aborted at all* the sites.
  - It's unacceptable to have a transaction committed at one site and aborted at another.
  - Consider the cross-site money transfer example again.
- The **two-phase commit** (2PC) protocol is widely used.

# Two-Phase Commit Protocol (2PC)

- Assumes **fail-stop** model
  - Failed sites simply stop working, and do not cause any other harm, such as sending incorrect messages to other sites.
- Execution of the protocol is initiated by the **Transaction Coordinator** after the *last step* of the transaction has been reached.
- The protocol involves all the sites at which the transaction is executed.
  - All participants
- Let  $T$  be a transaction initiated at site  $S_i$ , and let the transaction coordinator at  $S_i$  be  $TC_i$

# Phase 1: Obtaining a Decision

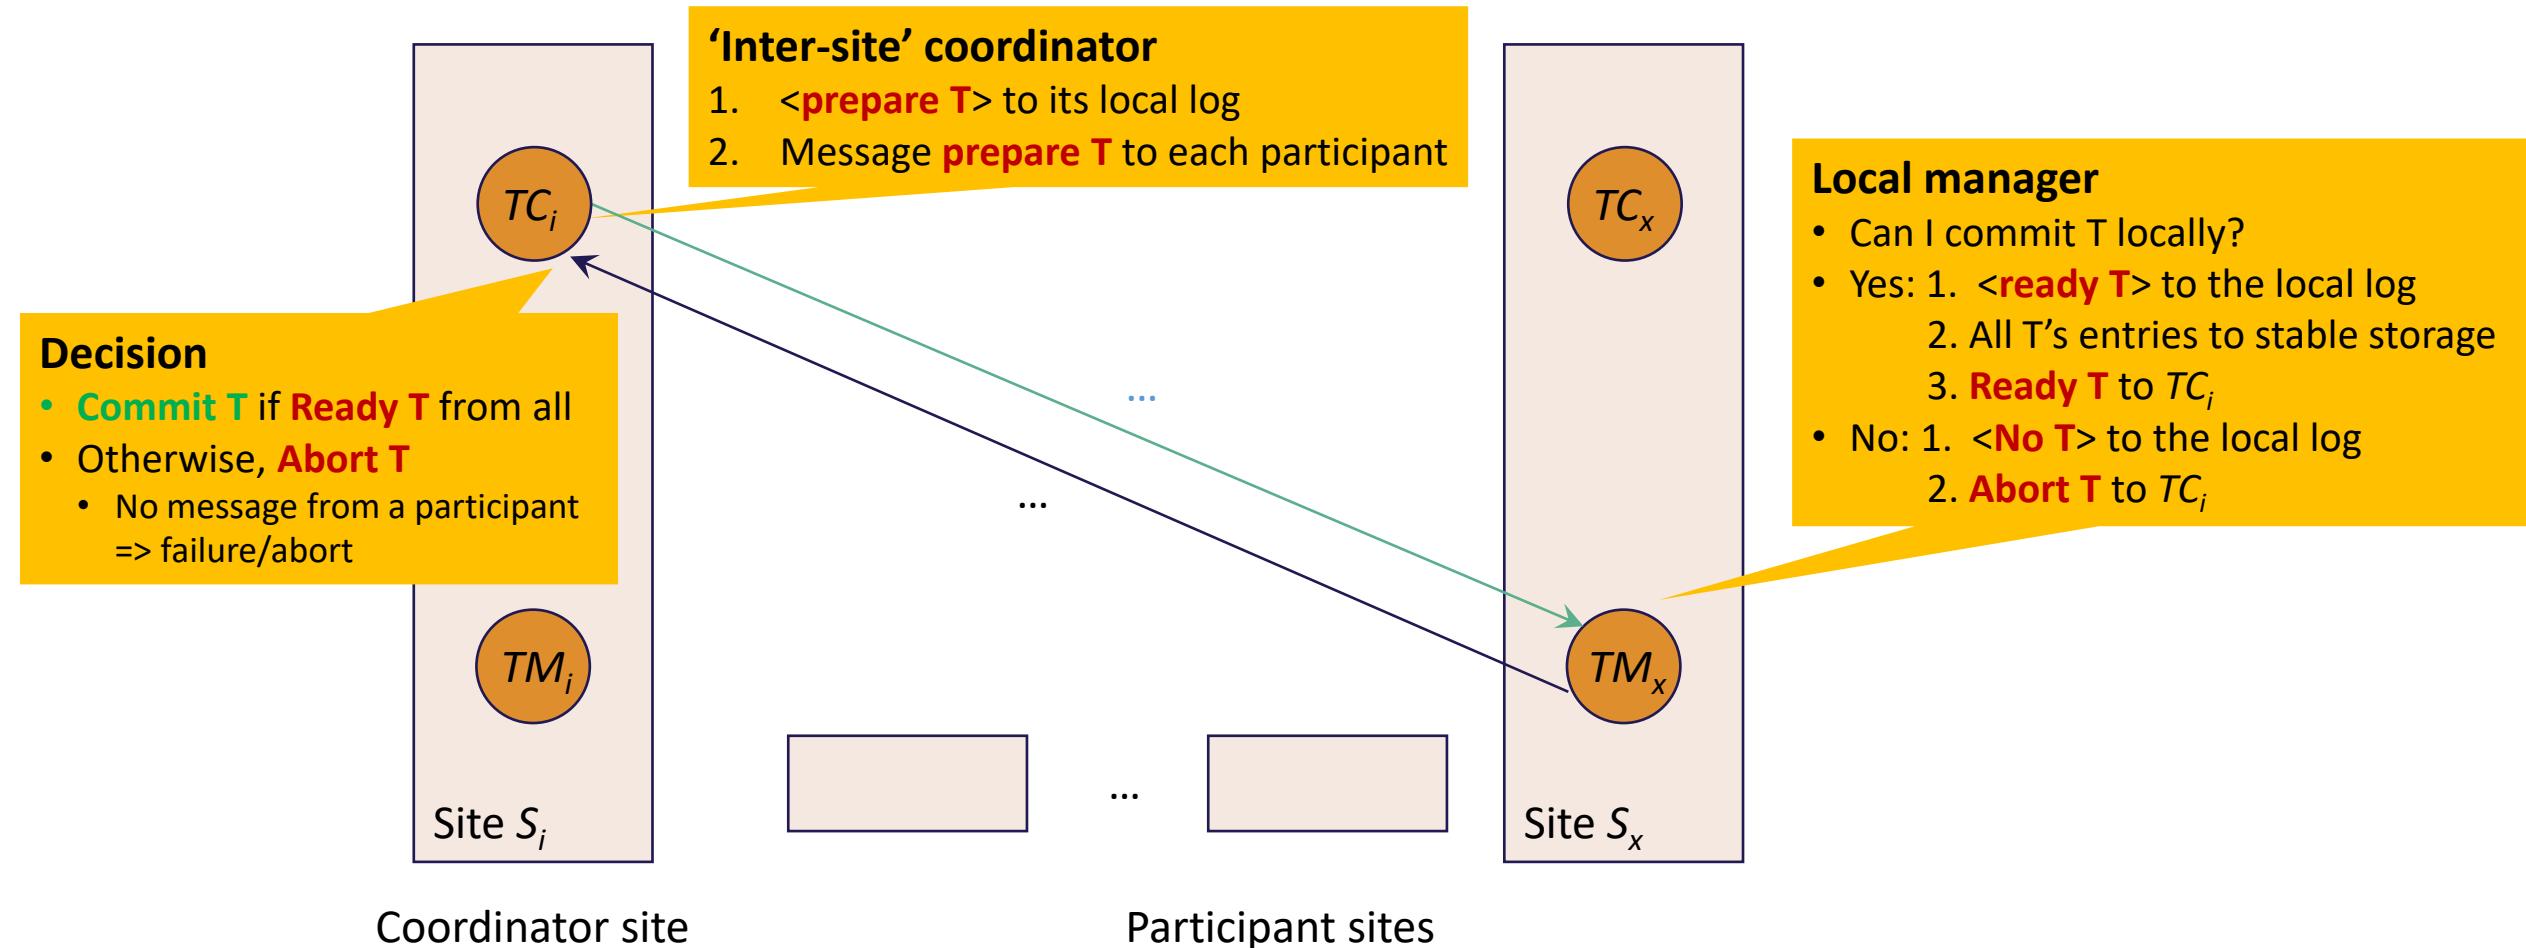
- ➊ Coordinator asks all participants to *prepare* to commit transaction  $T$ .
  - ➋  $TC_i$  adds the entry  $\langle \text{prepare } T \rangle$  to the log and forces log to stable storage.
  - ➋  $TC_i$  sends **prepare**  $T$  messages to all participants.
- ➋ Upon receiving the message, the **Transaction Manager** at a site determines if it can commit the transaction.
  - ➌ If not, add an entry  $\langle \text{no } T \rangle$  to the log and send **abort**  $T$  message to  $TC_i$ .
  - ➌ If the transaction can be committed, then:
    - › Adds the entry  $\langle \text{ready } T \rangle$  to the log
    - › Forces *all* entries for  $T$  to stable storage
    - › Send **ready**  $T$  message to  $TC_i$ .

# Phase 2: Recording the Decision

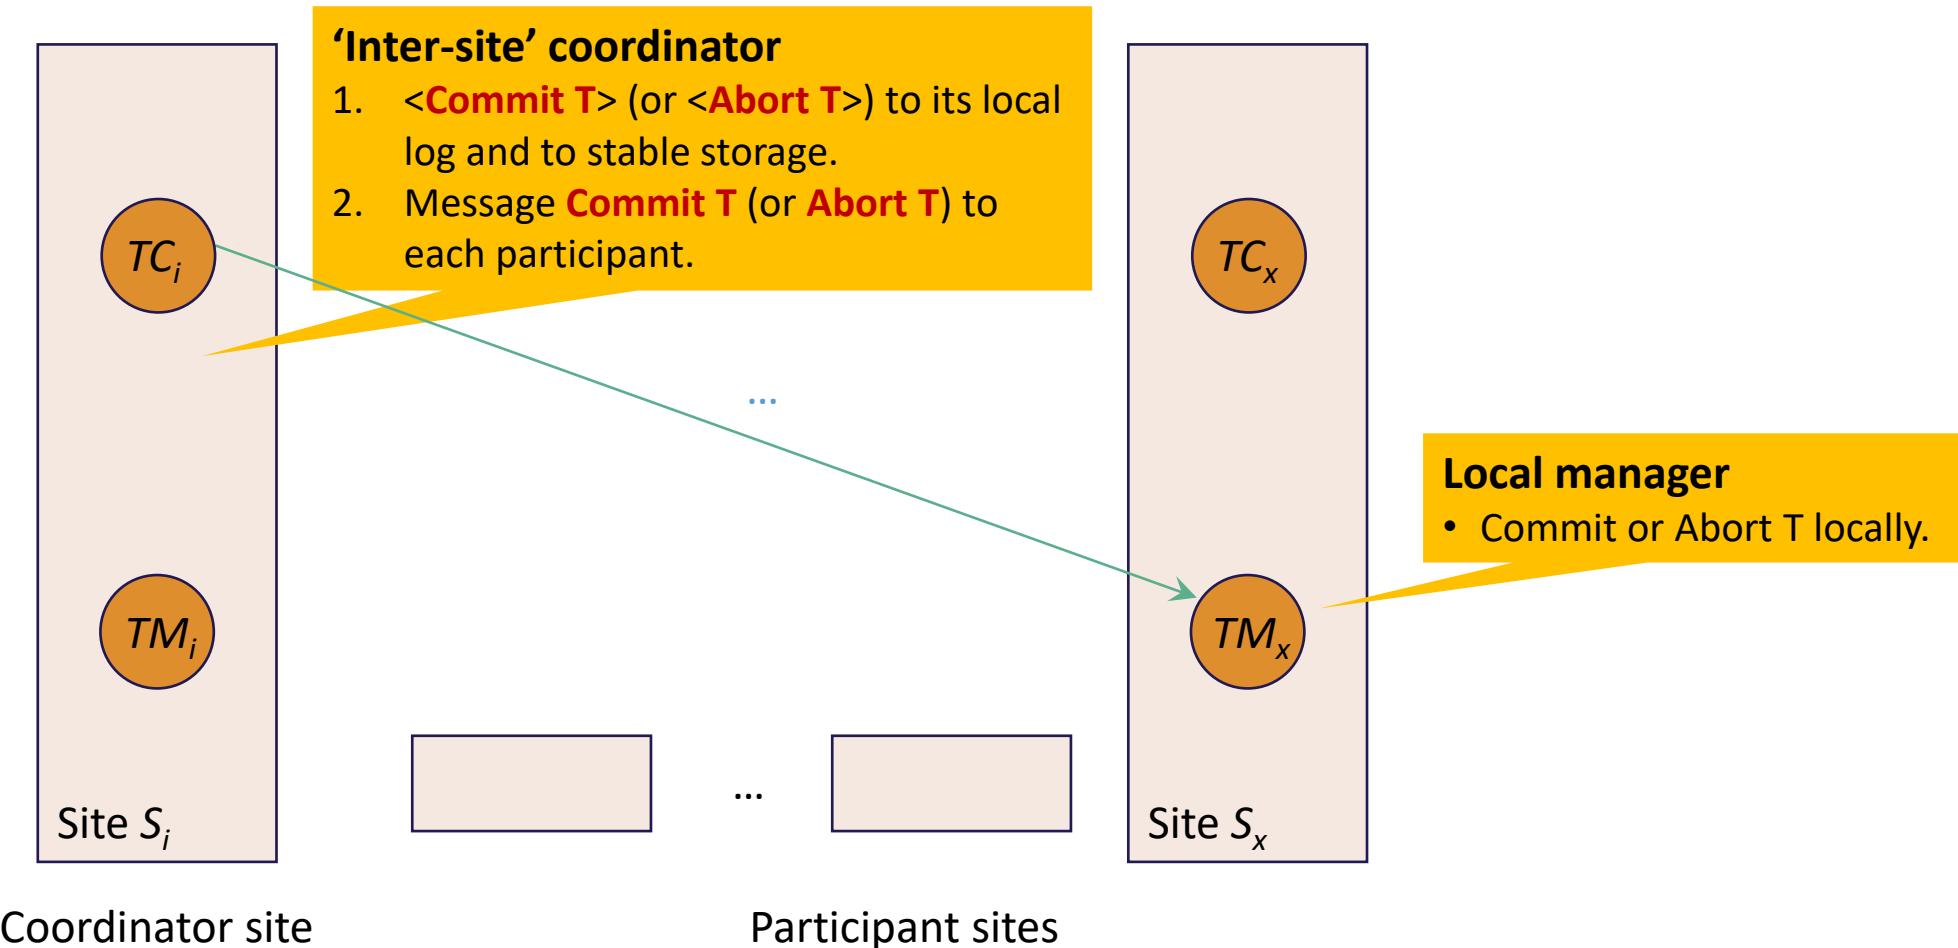
- $T$  can be committed if  $TC_i$  receives a **ready  $T$**  message from *all* the participating sites; otherwise,  $T$  must be aborted.
  - If no message from a participant, then also failure/abort.
- Coordinator adds a decision entry, **<commit  $T$ >** or **<abort  $T$ >**, to the log and forces the entry onto stable storage. Once the entry reaches stable storage it is irrevocable (even if failures occur).
- Coordinator sends a message to each participant, informing it of the decision (commit or abort).
- Each participant takes the appropriate action locally.

2PC is able to handle different types of failures.

# Illustration of Obtaining a Decision



# Illustration of Recording The Decision



# Handling of Participant Failure

- ➊ When site  $S_x$  recovers, it examines its log to determine the fate of transactions active at the time of the failure.
  - ➌ Log contain **<commit  $T$ >** entry: site executes **redo** ( $T$ )
  - ➌ Log contains **<abort  $T$ >** entry: site executes **undo** ( $T$ )
  - ➌ Log contains **<ready  $T$ >** entry: site must consult  $TC_i$  to determine the fate of  $T$ .
    - If  $T$  committed, **redo** ( $T$ )
    - If  $T$  aborted, **undo** ( $T$ )
- ➋ If the log contains no control entries concerning  $T$ ,  $S_x$  got failed before responding to the **prepare  $T$**  message from  $TC_i$ 
  - ➌ Since the failure of  $S_x$  precludes the sending of such a response,  $C_i$  must abort  $T$
  - ➌  $S_x$  must execute **undo** ( $T$ )

# Handling of Coordinator Failure

- If a coordinator fails while the commit protocol for  $T$  is executing, then participating sites must decide on  $T$ 's fate:
  - If an active site contains a **<commit  $T$ >** record in its log, then  $T$  must be committed.
  - If an active site contains an **<abort  $T$ >** record in its log, then  $T$  must be aborted.
  - If some active participating site does not contain a **<ready  $T$ >** record in its log, then the failed coordinator  $TC_i$  cannot have decided to commit  $T$ . Can therefore abort  $T$ .
  - If none of the above cases holds, then all active sites must have a **<ready  $T$ >** record in their logs, but no additional control records (such as **<abort  $T$ >** or **<commit  $T$ >**). In this case active sites must wait for  $TC_i$  to recover, to find decision.
- **Blocking problem:** active sites may have to wait for failed coordinator to recover.

# Handling of Network Partition

- If the coordinator and all its participants remain in one partition, the failure has no effect on the commit protocol.
- If the coordinator and its participants belong to several partitions:
  - Sites that are not in the coordinator's partition think the coordinator has failed and execute the protocol to deal with failure of the coordinator.
    - No harm results, but sites may still have to wait for decision from coordinator.
  - The coordinator and the sites in the same partition think that the sites in the other partition(s) have failed and follow the usual commit protocol.
    - Again, no harm results

# Summary

- Distributed Data Storage
  - Replication and fragmentation
- Distributed Query Processing
  - Cost measure
  - Query tree transformation in the presence of fragments
  - Join strategies, semijoin
- Distributed Transactions
  - Two phase commit protocol

# Readings

- Mandatory readings (for both Lectures 2 and 3)
  - A. Silberschatz, H. F. Korth, S. Sudarshan: Database System Concepts (7th edition), McGraw-Hill. Chapters 20, 21, 22 and 23
    - Optional: 21.3, 21.5, 22.6, 22.7, 22.8, 23.4, 23.5, 23.6, 23.7, 23.8

# Exercise 1

Consider the following setting and statistics

- $r_1$  with schema R(A, B) at site  $S_1$ ,  
 $\text{card}(R) = 10,000$   
 $\text{card}(\prod_A(r_1)) = 1,000$   
 $\text{card}(\prod_B(r_1)) = 2,000$   
 $\text{card}(r_1 \bowtie r_2) = 1,500$
- $r_2$  with schema S(A, C) at site  $S_2$   
 $\text{card}(S) = 50,000$   
 $\text{card}(\prod_A(r_2)) = 3,000$   
 $\text{card}(\prod_C(r_2)) = 5,000$   
 $\text{card}(r_2 \bowtie r_1) = 3,500$
- Attribute data value size:
  - A: 10 bytes; B: 20 bytes; C: 30 bytes.
- Basic costs:
  - $c_0$ : initial set-up cost between any two sites = 20
  - $c_1$ : transmission cost per data unit between any two sites = 1 per 500 bytes
- **Questions:** If we use the semijoin strategy to obtain  $r_1 \bowtie r_2$ ,
  - What is the total cost to obtain the result at  $S_1$ ?
  - What is the total cost to obtain the result at  $S_2$ ?

# Exercise 2

## Site 1: **EMPLOYEE**

- Schema: (EID, Name, Salary, DID)
  - EID: 10 bytes
  - Name: 20 bytes
  - Salary: 20 bytes
  - DID: 10 bytes
  - Totally 1000 tuples

## Site 2: **DEPARTMENT**

- Schema: (DID, DName)
  - DID: 10 types
  - Dname: 20 types
  - Totally 50 tuples

- **Question:** Site 3 needs to find the name of employees and their department names. Figure out at least 3 strategies for this distributed join query and calculate the total amount of data transfer for each strategy.